

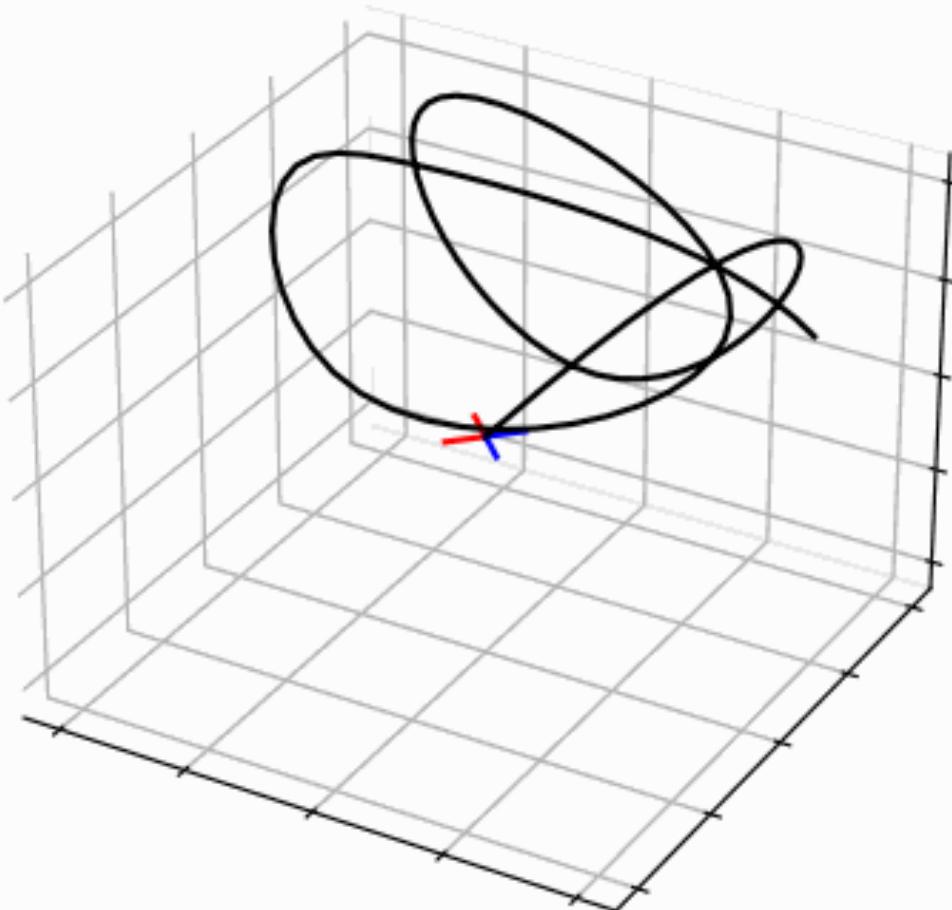
# **Learning Algorithm Hyperparameters for Fast Parametric Convex Optimization with Certified Robustness**

**ICCOPT 2025**  
**Rajiv Sambharya**

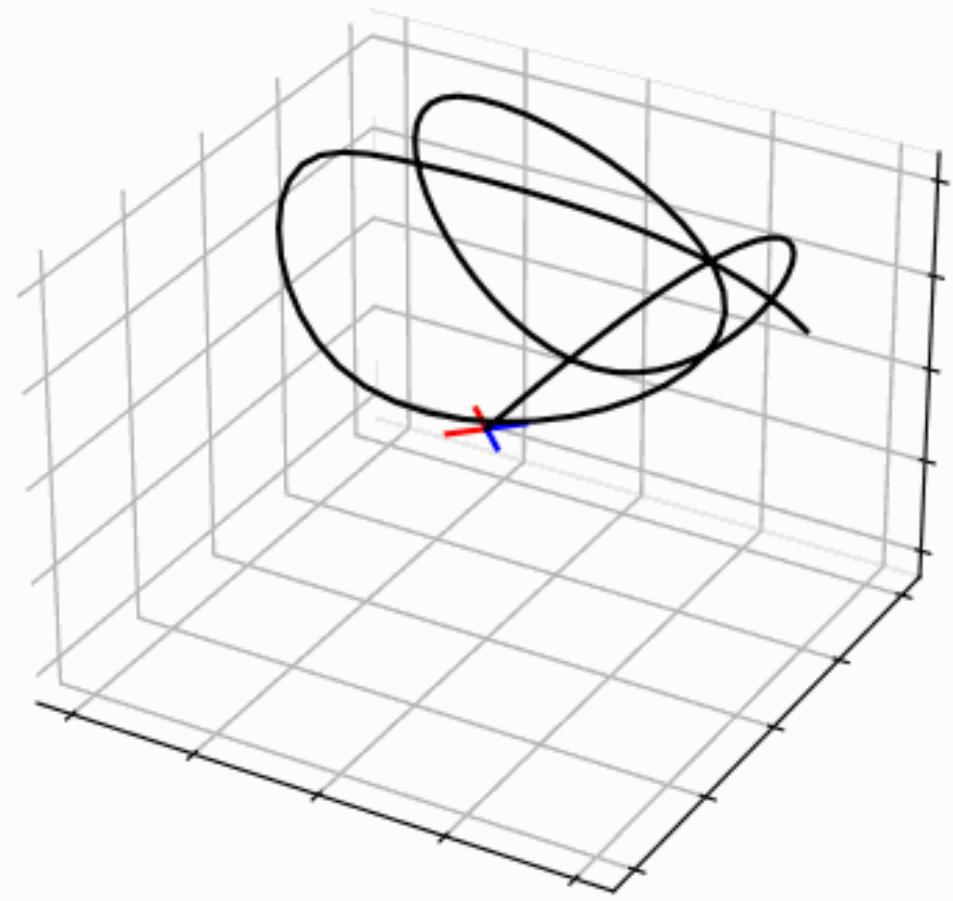


**Penn**  
UNIVERSITY of PENNSYLVANIA

# Tracking a reference trajectory with a quadcopter



Success!  
(If given enough time)



Failure: not enough time to solve

## Model predictive control

optimize over a smaller horizon ( $T$  steps),  
implement first control,  
repeat

### Model predictive controller

$$\begin{aligned} & \text{minimize} && \sum_{t=1}^T \|s_t - s_t^{\text{ref}}\|_2^2 \\ & \text{subject to} && s_{t+1} = As_t + Bu_t \\ & && s_t \in \mathcal{S}, \quad u_t \in \mathcal{U} \\ & && s_0 = s_{\text{init}} \end{aligned}$$

Current state,  
reference trajectory

Control  
inputs

# Real-world optimization is parametric

Parameter

$$x \longrightarrow$$

minimize  $f(z, x)$   
subject to  $g(z, x) \leq 0$   
 $f$  and  $g$  convex in  $z$

Optimal solution

$$\longrightarrow z^*(x)$$

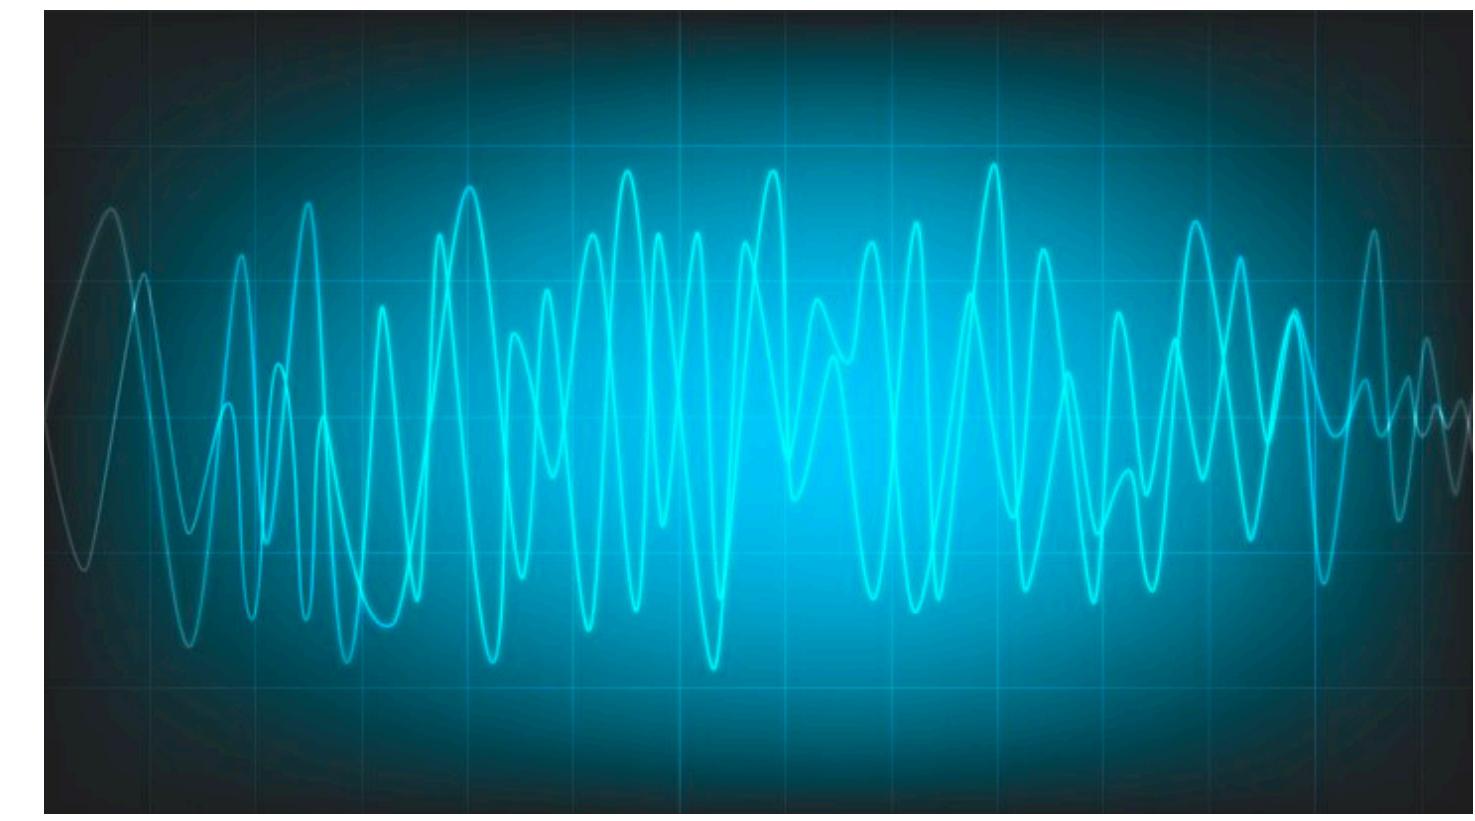
Robotics and control



Energy



Signal processing



# First-order methods are widely popular again...

First-order methods use only gradient information

Fixed-point iterations  $z^{k+1}(x) = T(z^k(x), x)$

## Example: projected gradient descent

minimize  $\mathbf{f}(z, x)$  convex  
smooth

subject to  $z \in \mathcal{C}(x)$  convex set

$$z^{k+1}(x) = \Pi_{\mathcal{C}(x)} \left( z^k(x) - \underbrace{\alpha \nabla \mathbf{f}(z^k(x), x)}_{\text{gradient step}} \right)$$

projection

## Benefits of first-order methods

cheap iterations 

embedded optimization



large-scale optimization



Many general-purpose solvers...

PDLP   
Applegate et al. 2021

OSQP   
Stellato et al. 2020

SCS   
SPLITTING CONIC SOLVER  
O'Donoghue et al. 2016

# ...But general-purpose first-order methods can converge slowly

Initialize

$$z^0(x) = 0$$

Algorithm steps

$$z^{k+1}(x) = T(z^k(x), x)$$

Terminate when

$$\|z^{k+1}(x) - z^k(x)\|_2 \leq \epsilon$$

Fixed point

$$z^*(x) = T(z^*(x), x)$$

Optimal solution



Problem!

In many applications, we have a **budget** of iterations  
(e.g., I only have the time to run 20 fixed-point steps)

# Can machine learning speed up convex parametric optimization?

**Goal: Do mapping quickly and accurately**

Parameter

$$x \longrightarrow$$

$$\begin{aligned} & \text{minimize} && f(z, x) \\ & \text{subject to} && g(z, x) \leq 0 \end{aligned}$$

Optimal solution

$$\longrightarrow z^*(x)$$

Only Optimization



$$\longrightarrow \hat{z}^{\text{Opt}}(x)$$



Only Machine Learning



$$\longrightarrow \hat{z}^{\text{ML}}(x)$$



Optimization Machine Learning



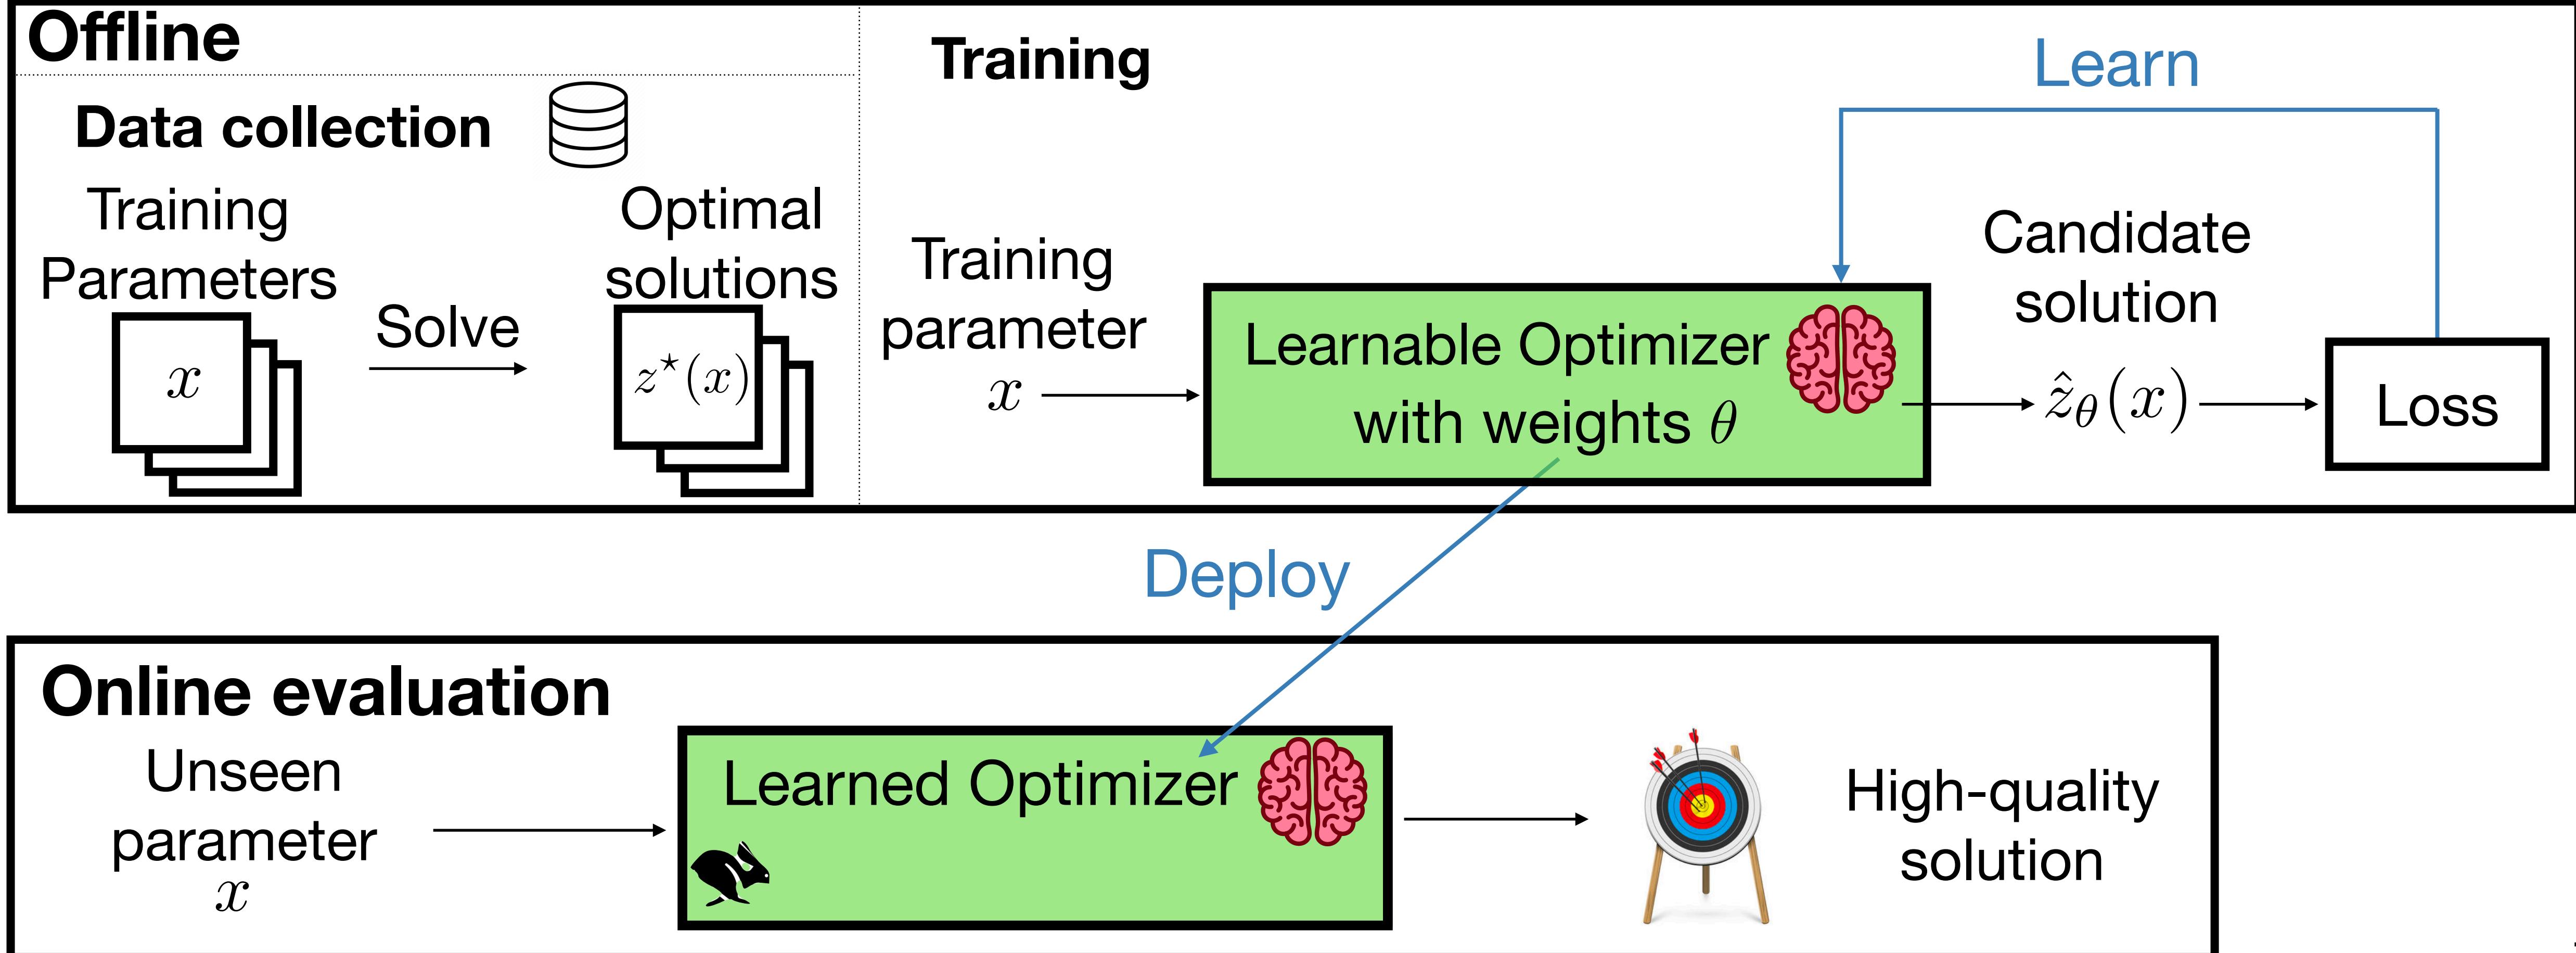
$$\longrightarrow \hat{z}^{\text{Opt+ML}}(x)$$



# The learning to optimize paradigm

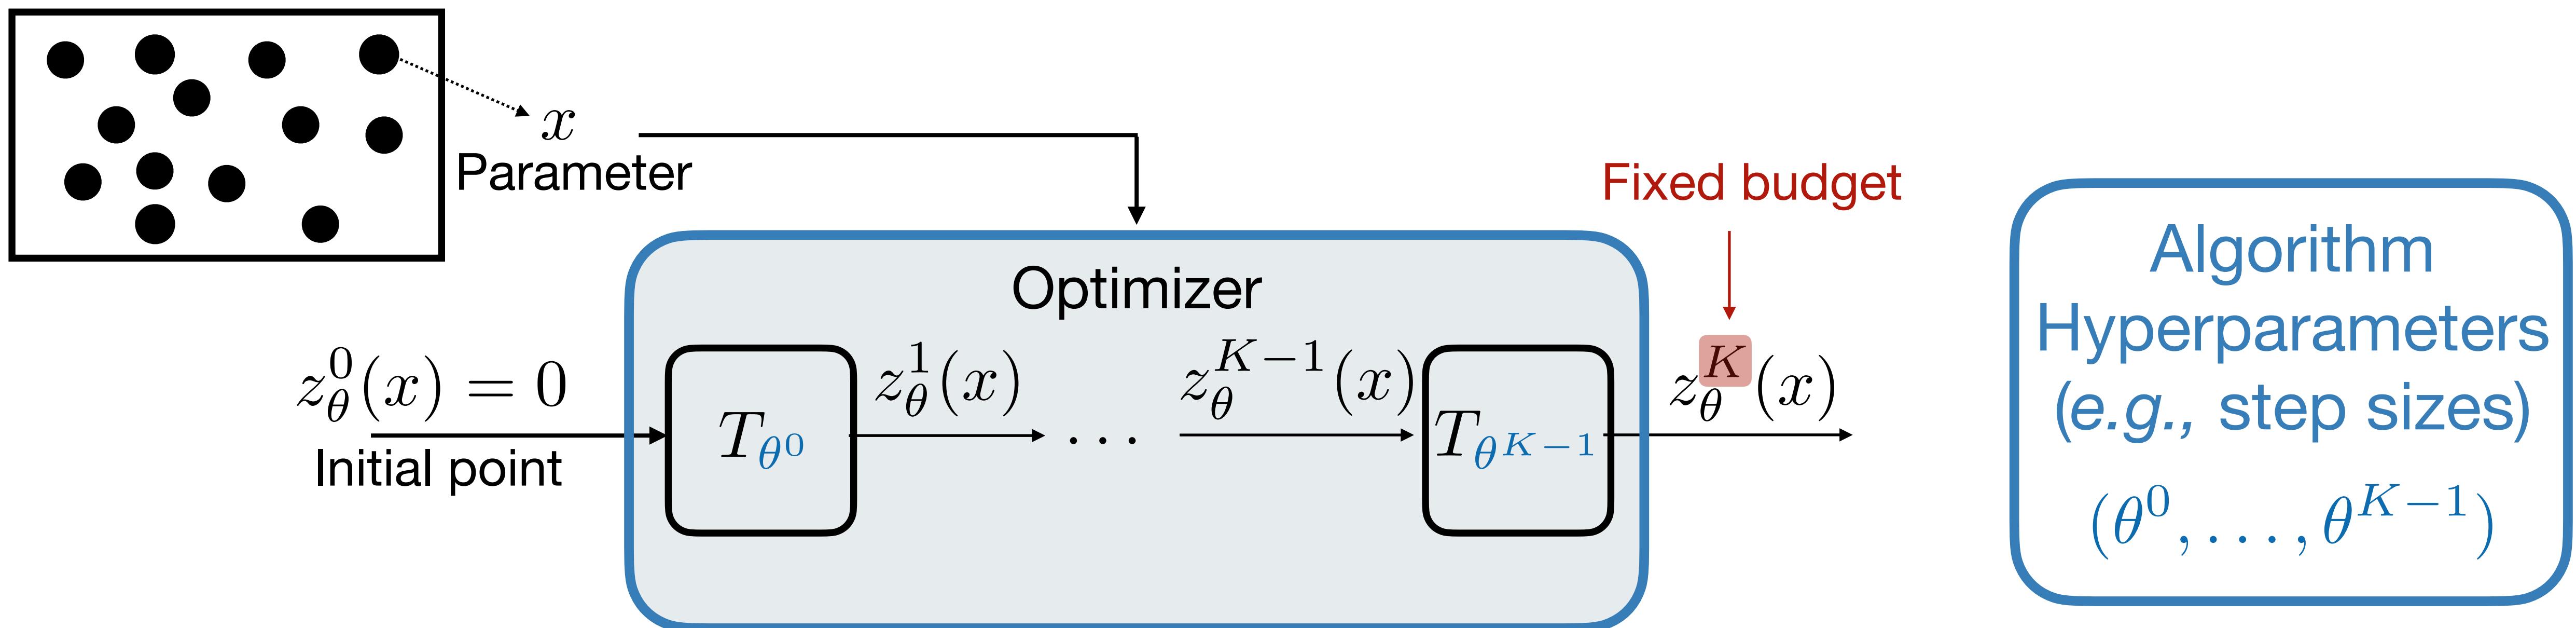
**Goal:** solve the parametric optimization problem fast

$$\begin{aligned} & \text{minimize} && f(z, x) \\ & \text{subject to} && g(z, x) \leq 0 \end{aligned}$$



# Learning Algorithm Hyperparameters

# First-order methods as fixed-length computational graphs



Example: projected gradient descent

$$z_\theta^{k+1}(x) = \Pi_{\mathcal{C}(x)}(z_\theta^k(x) - \theta^k \nabla_z f(z_\theta^k(x), x))$$

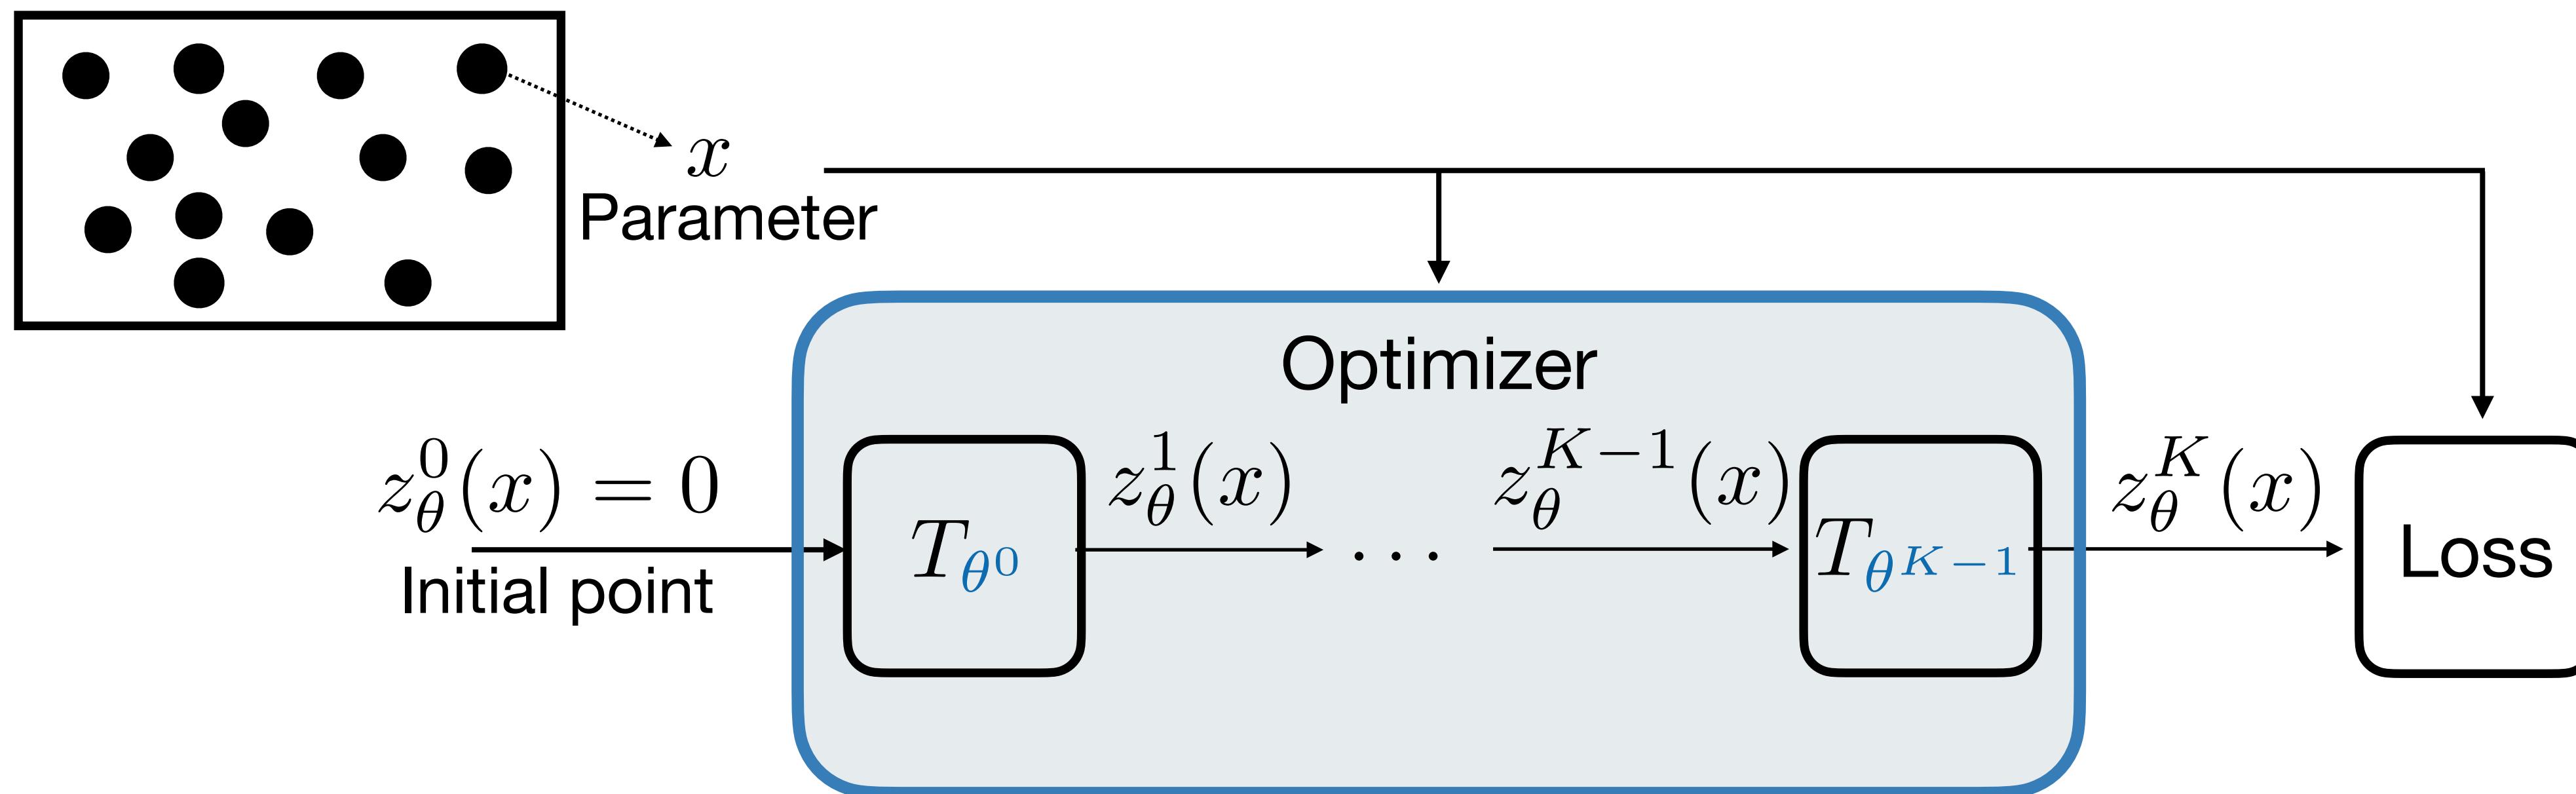
- Conventional wisdom: use a constant step size
- Recent advances: vary the step size!

Altschuler et. al 2023, Grimmer 2023, Bok et. al 2024



What if we **learn** the step sizes?

# Learning the algorithm hyperparameters framework



Provided  $N$   
training instances:  $\{(x_i, z^*(x_i))\}_{i=1}^N$

Training problem  
 $\begin{aligned} \text{minimize } & (1/N) \sum_{i=1}^N \|z_\theta^K(x_i) - z^*(x_i)\|_2^2 \\ \text{subject to } & z_\theta^{k+1}(x_i) = T_{\theta^k}(z_\theta^k(x_i)) \\ & z_\theta^0(x_i) = 0 \end{aligned}$

Optimize  $\theta$  with  
gradient-based methods

Learned hyperparameters  

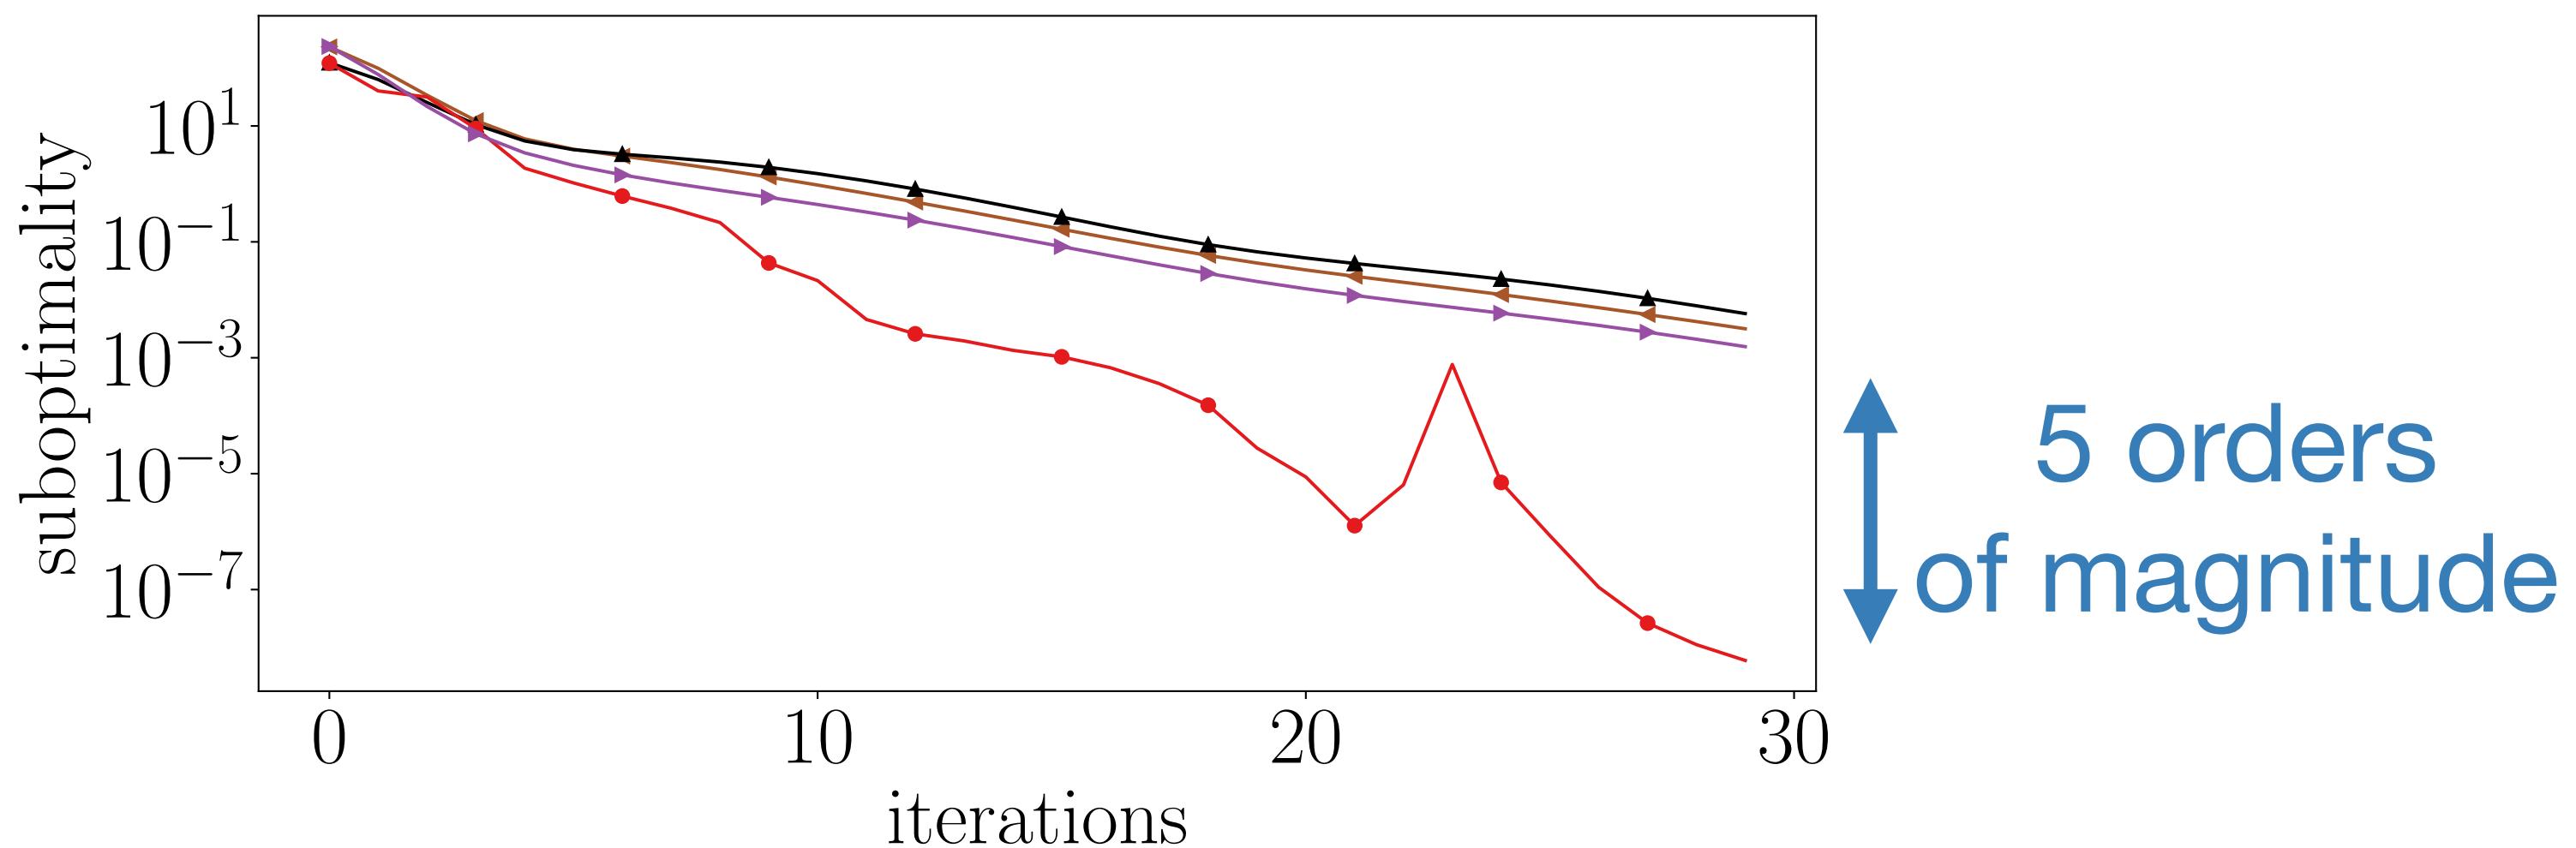
- shared across problem instances
- differ across iterations

# Learning step sizes for non-negative least squares

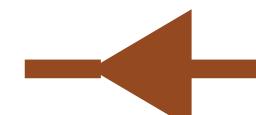
Parameter  
 $x$

$$\begin{aligned} & \text{minimize} && (1/2)\|Az - x\|_2^2 \\ & \text{subject to} && z \geq 0 \end{aligned}$$

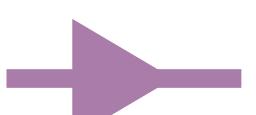
Solution  
 $z^*(x)$



Nearest neighbor  
warm start



Learned  
warm starts



Sambharya et. al 2024

10000 training instances

Learned  
step sizes



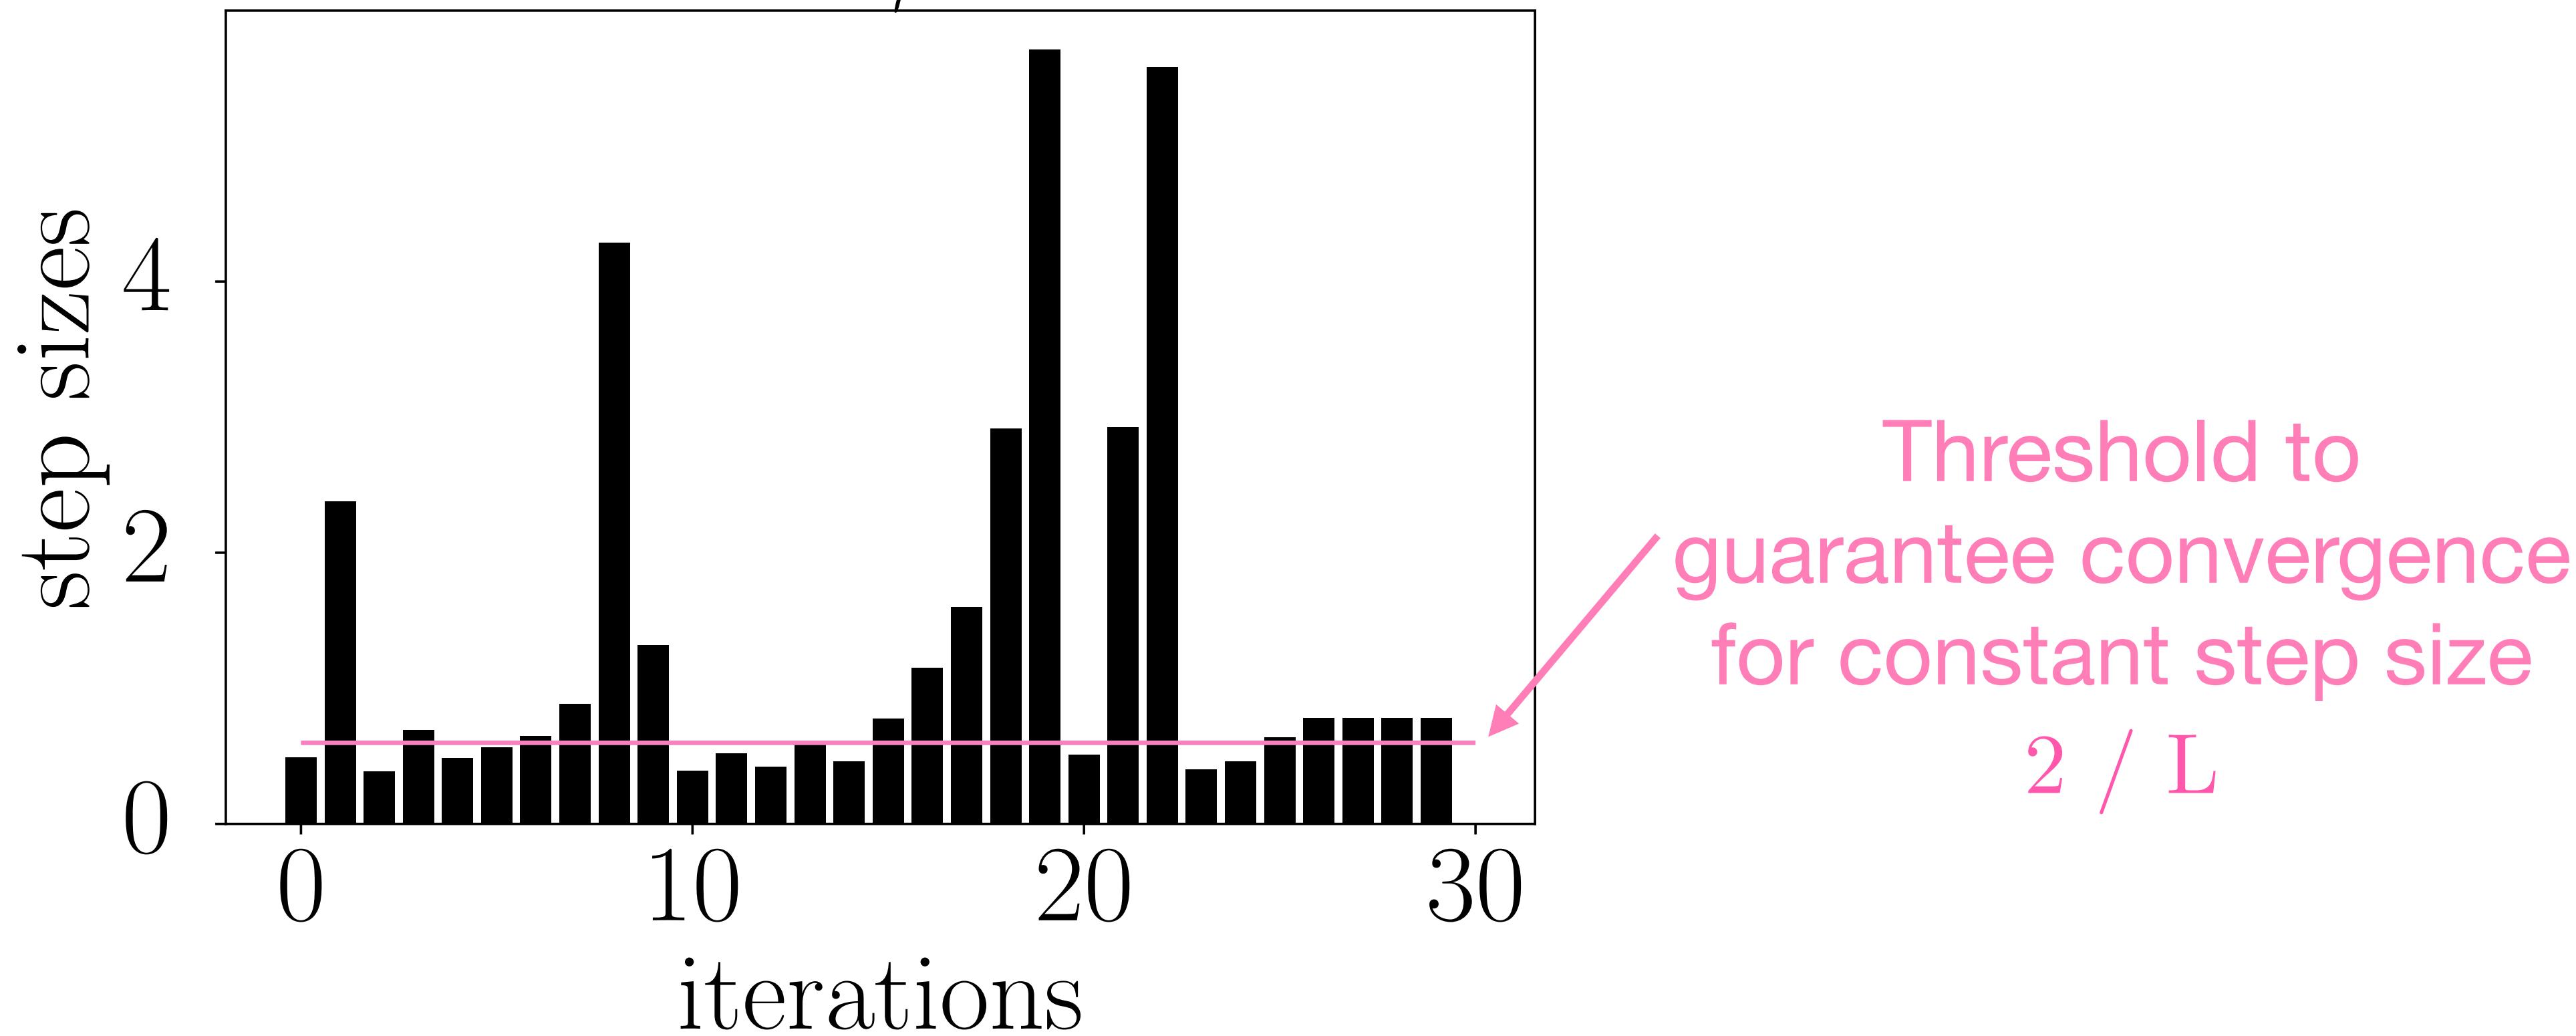
10 training instances

Learning step sizes  
can be powerful

This is a highly  
data-efficient approach

Multi-step  
descent phenomenon

# We learn long steps!



# An extension: we can also learn momentum sizes

## Composite convex optimization

$$\text{minimize } f(z, x) + g(z, x)$$

Convex, smooth      Convex

Proximal operator

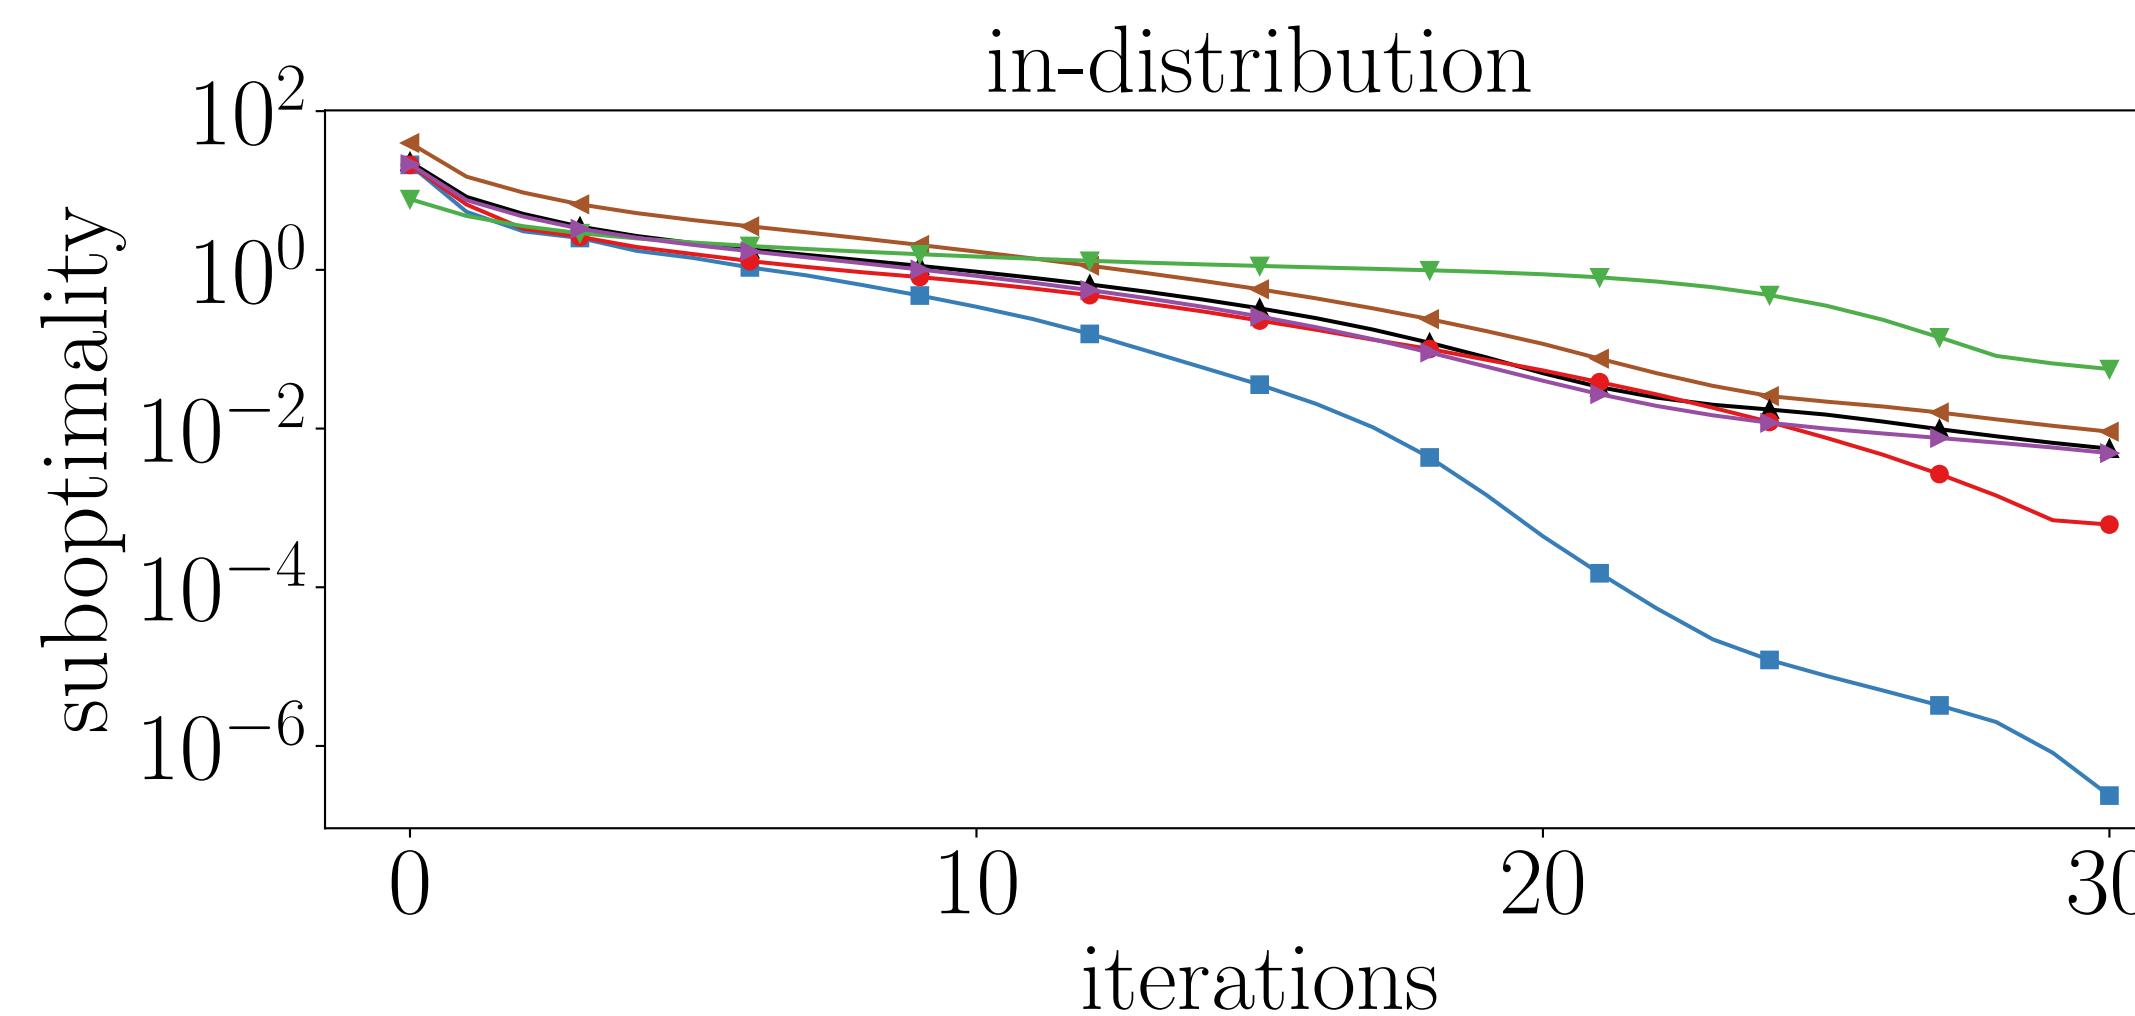
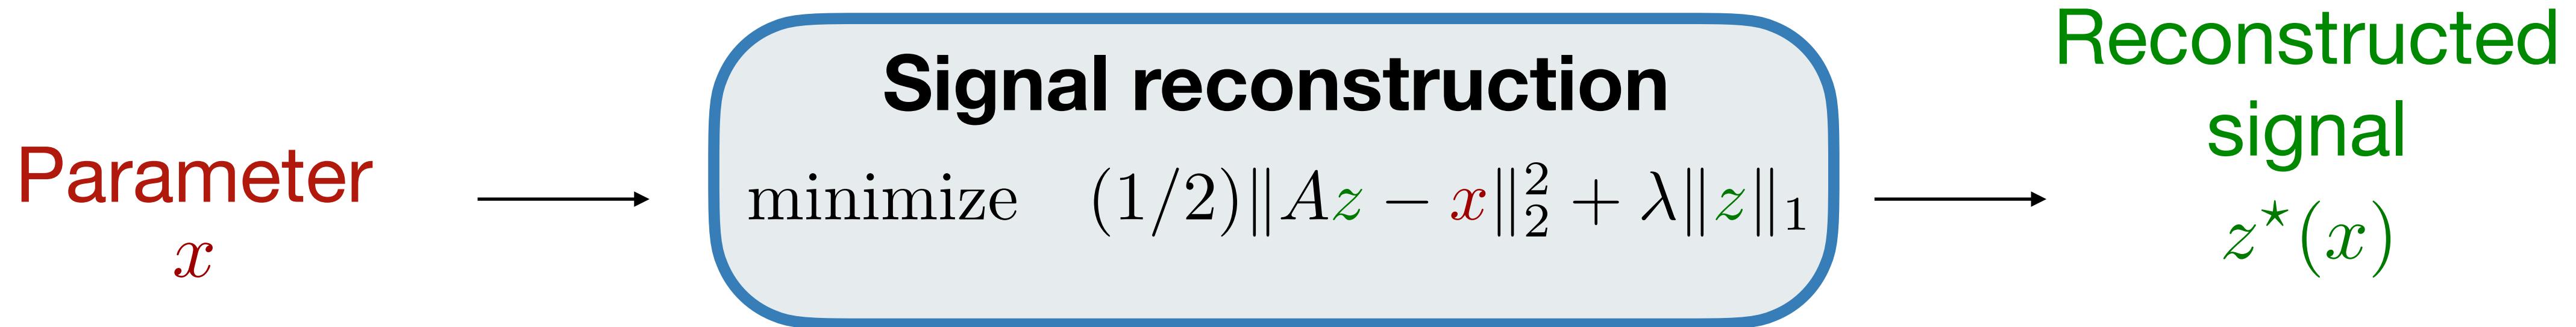
$$\mathbf{prox}_g(v) = \arg \min_u g(u) + (1/2)\|u - v\|_2^2$$

### Nesterov's acceleration

$$y^{k+1}(x) = \mathbf{prox}_{\alpha^k g}(z^k(x) - \alpha^k \nabla f(z^k(x), x))$$
$$z^{k+1}(x) = y^{k+1}(x) + \beta^k (y^{k+1}(x) - y^k(x))$$

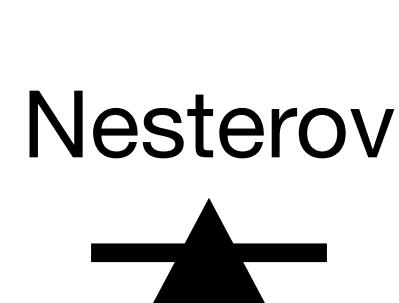
Learn  $\theta^k = (\alpha^k, \beta^k)$

# Example: learned hyperparameters for sparse coding

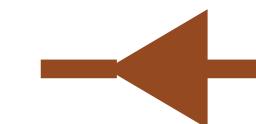


Learning momentum steps can sometimes help significantly

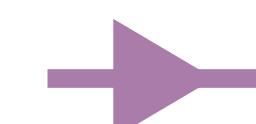
But what about worst-case guarantees?



Nearest neighbor  
warm start



Learned  
warm starts



Sambharya et. al 2024



Learned  
step sizes



Learned step and  
momentum sizes



10000 training instances

Gregor and LeCun 2010

10 training instances

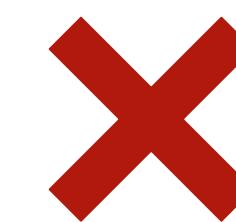
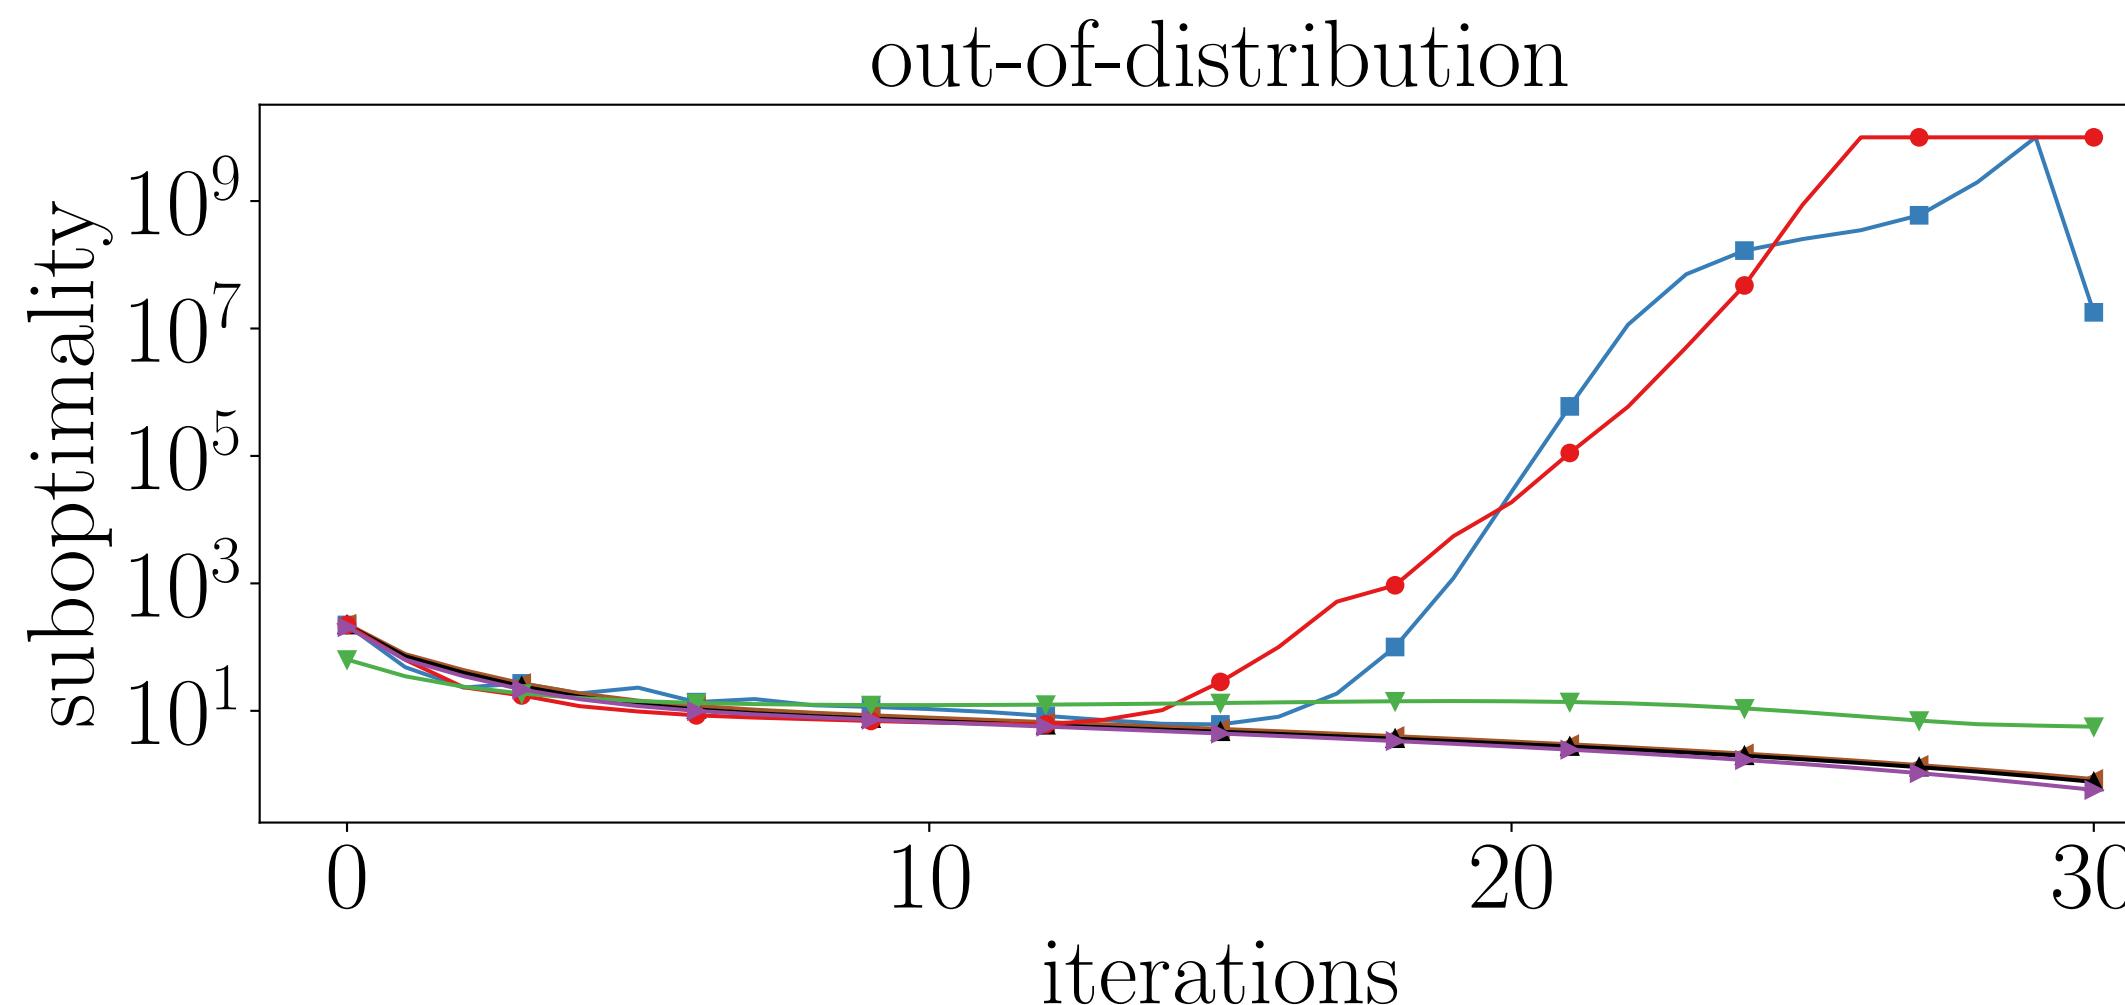
# This approach lacks worst-case guarantees

Parameter  
 $x$

## Signal reconstruction

$$\text{minimize} \quad (1/2)\|A\mathbf{z} - \mathbf{x}\|_2^2 + \lambda\|\mathbf{z}\|_1$$

Reconstructed  
signal  
 $\mathbf{z}^*(x)$



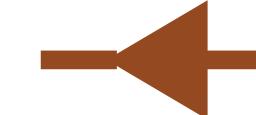
Failures on  
out-of-distribution instances



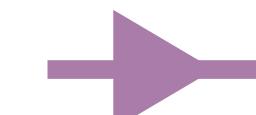
Can we learn hyperparameters  
that are robust?



Nearest neighbor  
warm start



Learned  
warm starts



Learned  
step sizes



Learned step and  
momentum sizes



# Certifying robustness of algorithms with learned hyperparameters

# Can we learn hyperparameters that are robust?

A strong form of robustness—worst-case guarantees for all parameters in a set  $\mathcal{X}$

$$r(z_\theta^K(x), x) \leq \gamma(\theta) \|z^0(x) - z^*(x)\|^2 \quad \forall x \in \mathcal{X}$$

Performance metric: e.g.,

$$r(z_\theta^K(x), x) = \|z_\theta^K(x) - z^*(x)\|_2^2$$

$$r(z_\theta^K(x), x) = f(z_\theta^K(x), x) - f(z^*(x), x)$$

Level of robustness

A provided set of interest

Ideally, learn  $\theta$  as before  
but constrain  $\gamma(\theta) \leq \gamma^{\text{target}}$



But how can we  
evaluate  $\gamma(\theta)$ ?

# Certified robustness for all parameters in a set

Definition:  $(f, \mathcal{X})$  is  $\mathcal{F}_{\mu, L}$ -parametrized if  $f(\cdot, x) \in \mathcal{F}_{\mu, L} \quad \forall x \in \mathcal{X}$

$\mu$ -strongly convex,  $L$ -smooth

Example: minimize  $\underbrace{(1/2)\|Az - x\|_2^2}_{f(z, x)} + \underbrace{\lambda\|z\|_1}_{g(z, x)} \quad \mathcal{X} = \mathbf{R}^d$

$(f, \mathcal{X})$  is  $\mathcal{F}_{\mu, L}$ -parametrized

min and max eigenvalues of  $A^T A$

$(g, \mathcal{X})$  is  $\mathcal{F}_{0, \infty}$ -parametrized



Worst-case guarantees over function class imply worst-case guarantees over set

Can leverage Performance Estimation Problem Analysis

# The Performance Estimation Problem (PEP) Framework can help us

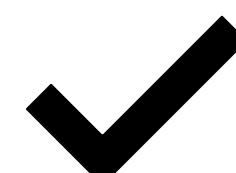
maximize	(performance metric)	$r(z^K)$
subject to	(initial point)	$z^0 = y^0, \ z^0 - z^*\ _2^2 \leq 1$
	(optimality)	$\nabla f(z^*) + \partial g(z^*) = 0$
	(algorithm update)	$y^{k+1} = \text{prox}_{\alpha^k g}(z^k - \alpha^k \nabla f(z^k))$
	(function class)	$z^{k+1} = y^{k+1} + \beta^k (y^{k+1} - y^k)$
		$f \in \mathcal{F}_{\mu,L}, g \in \mathcal{F}_{0,\infty}$ .

PEP: Tight convex SDP  
formulation using gram matrix  $G$

Drori, Teboulle, Hendrickx, Glineur, Taylor, Ryu, Grimmer, and many more

$$\begin{aligned} & \theta \longrightarrow \text{maximize} && \text{tr}(A_0 G) \\ & \text{subject to} && \text{tr}(A_i(\theta)G) \leq b_i, i = 1, \dots, m \\ & && G \succeq 0 \end{aligned}$$

Level of robustness  $\gamma(\theta)$



Solve an SDP to evaluate  $\gamma(\theta)$



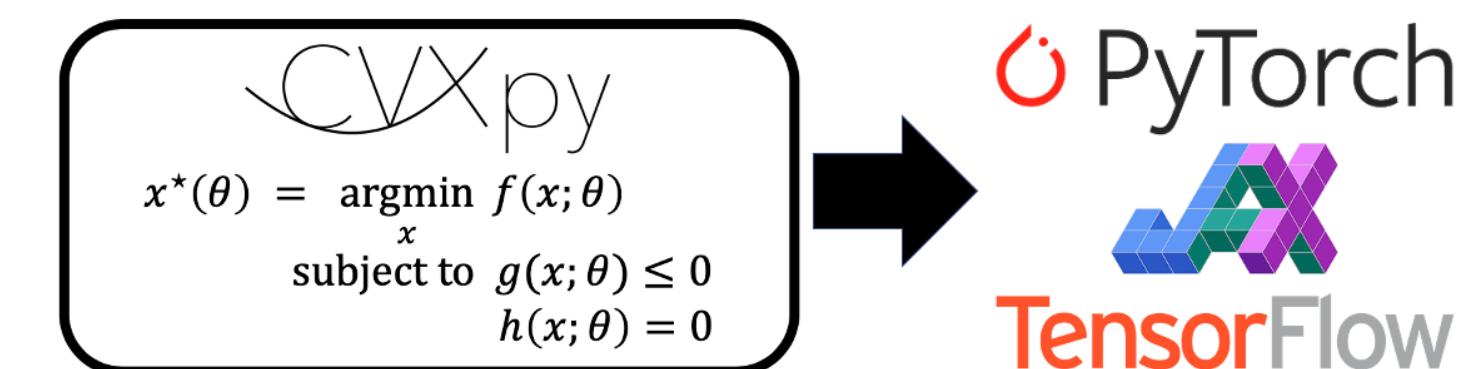
But how can we train for robustness?

# Robust training of hyperparameters

## PEP-regularized training problem

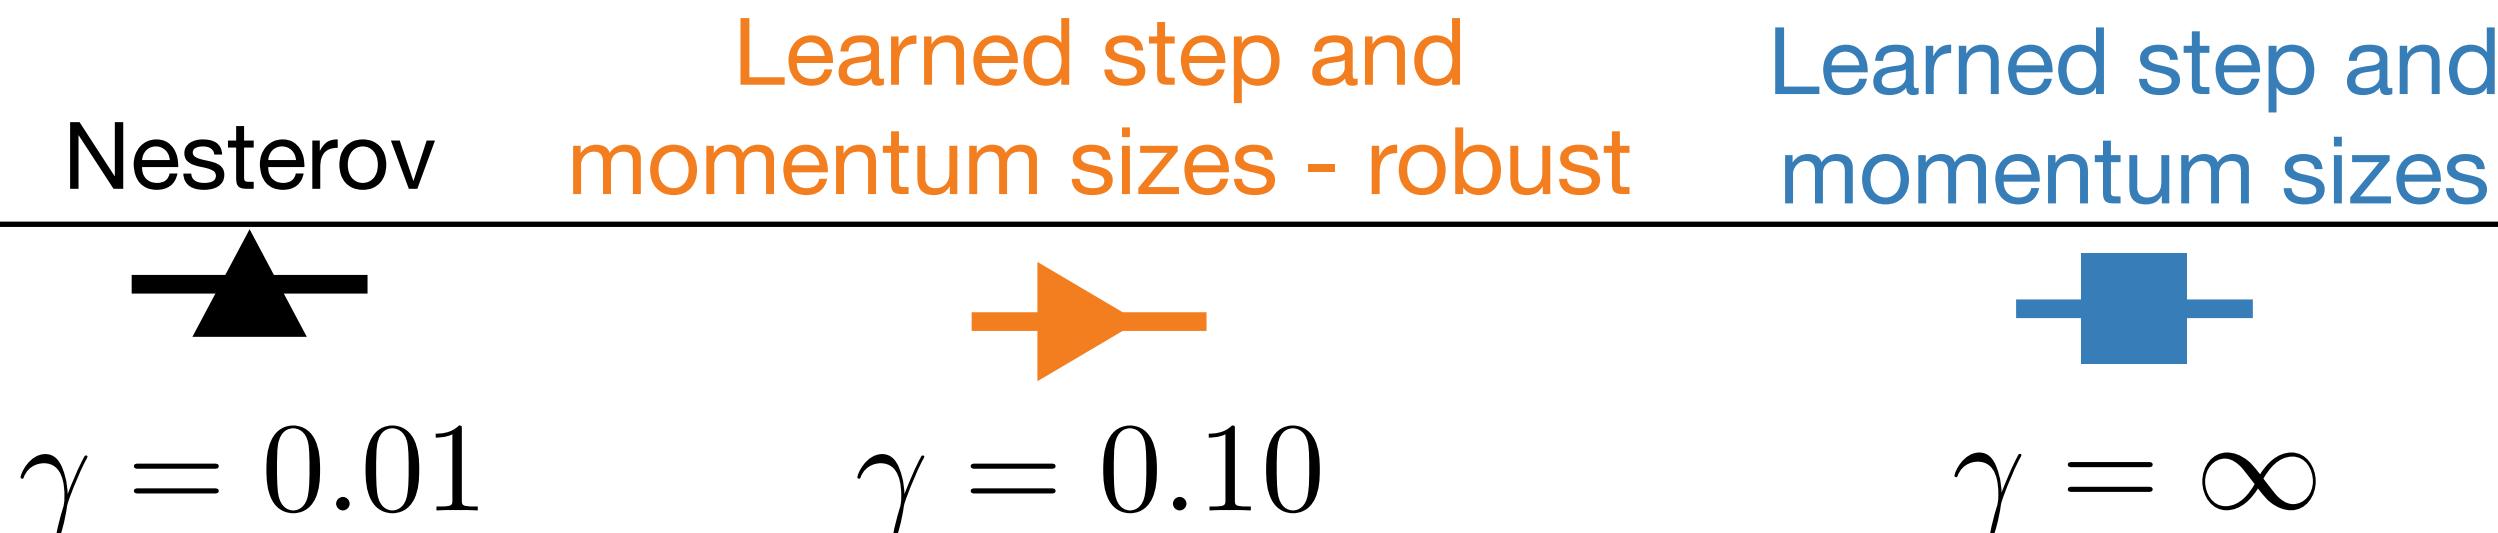
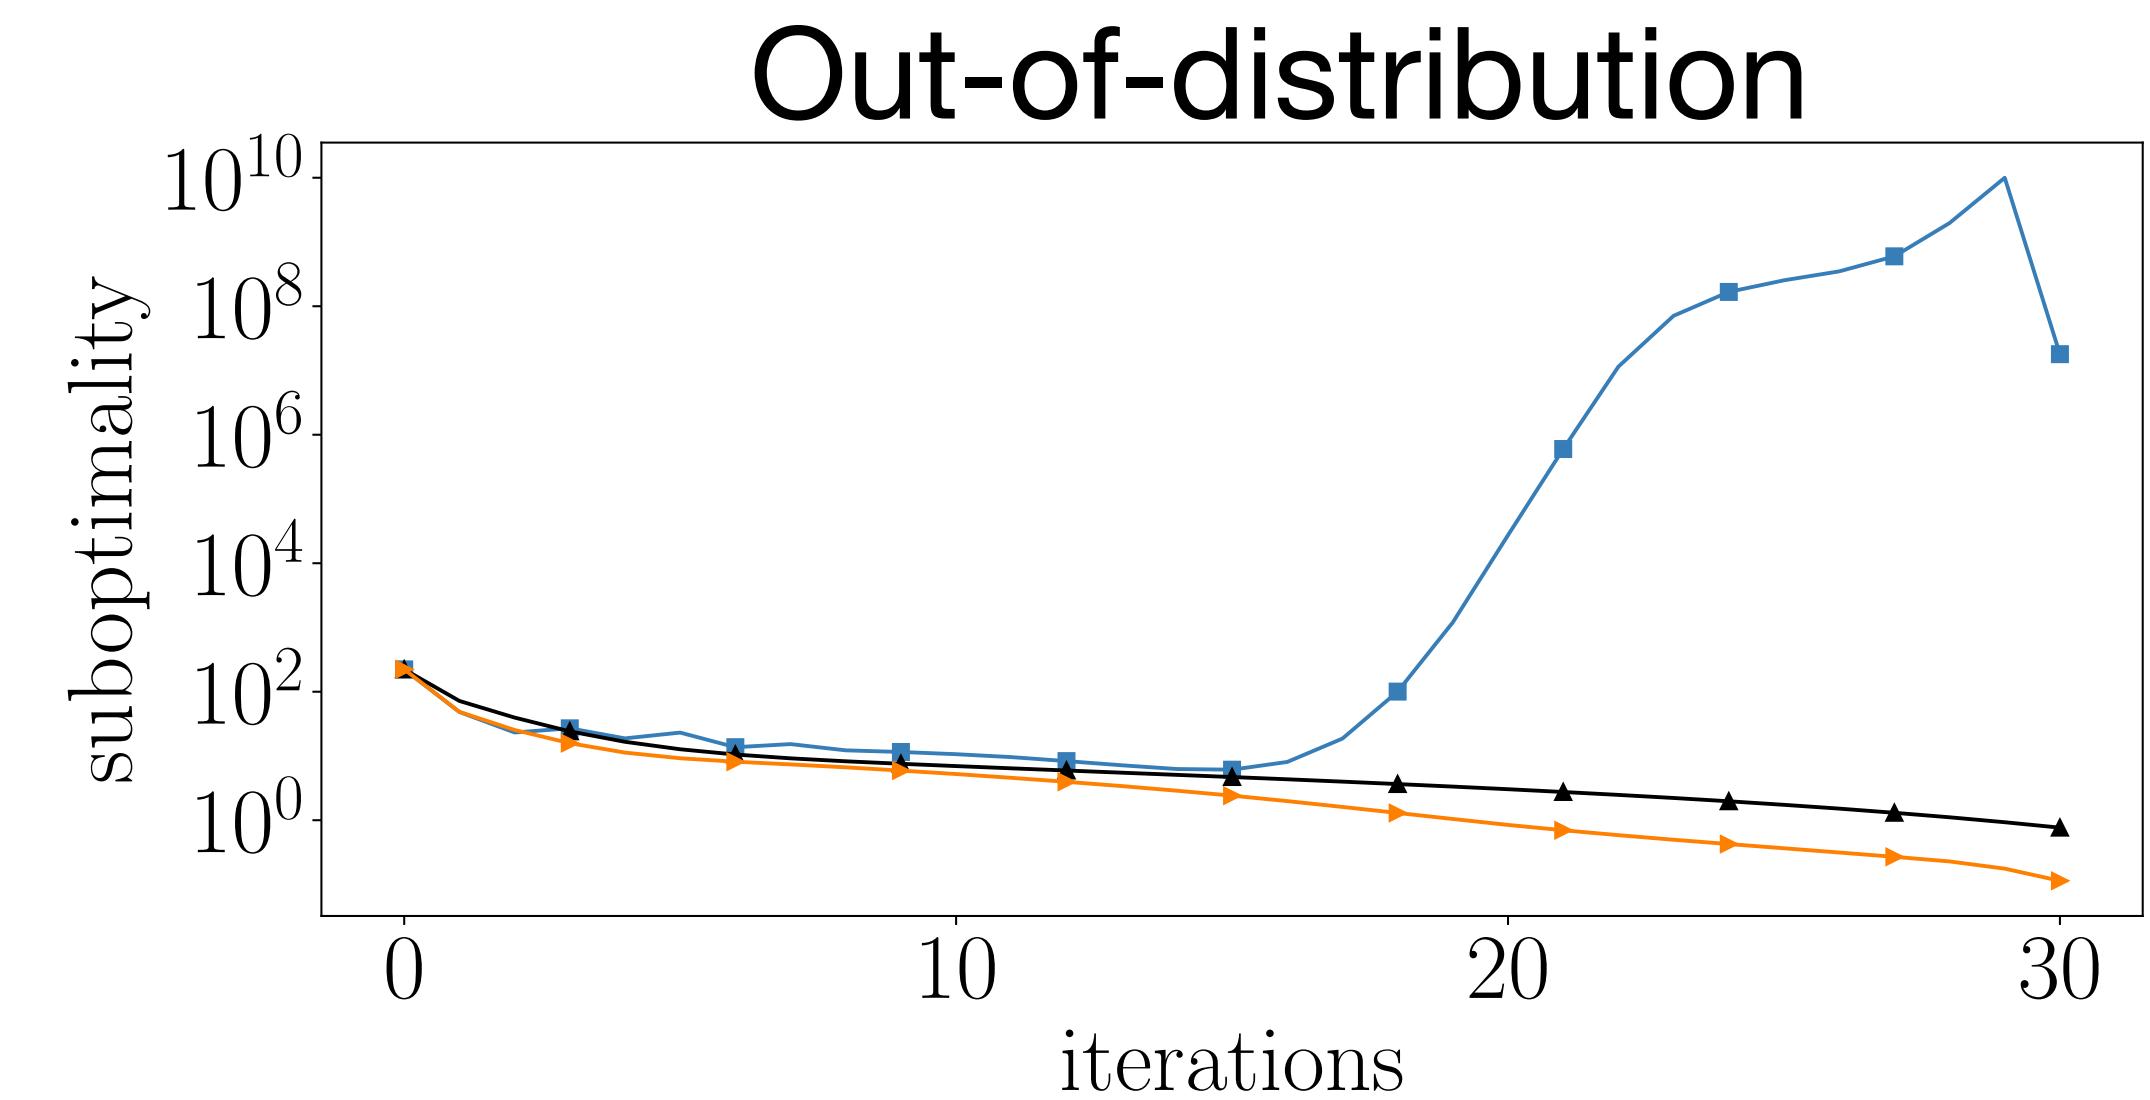
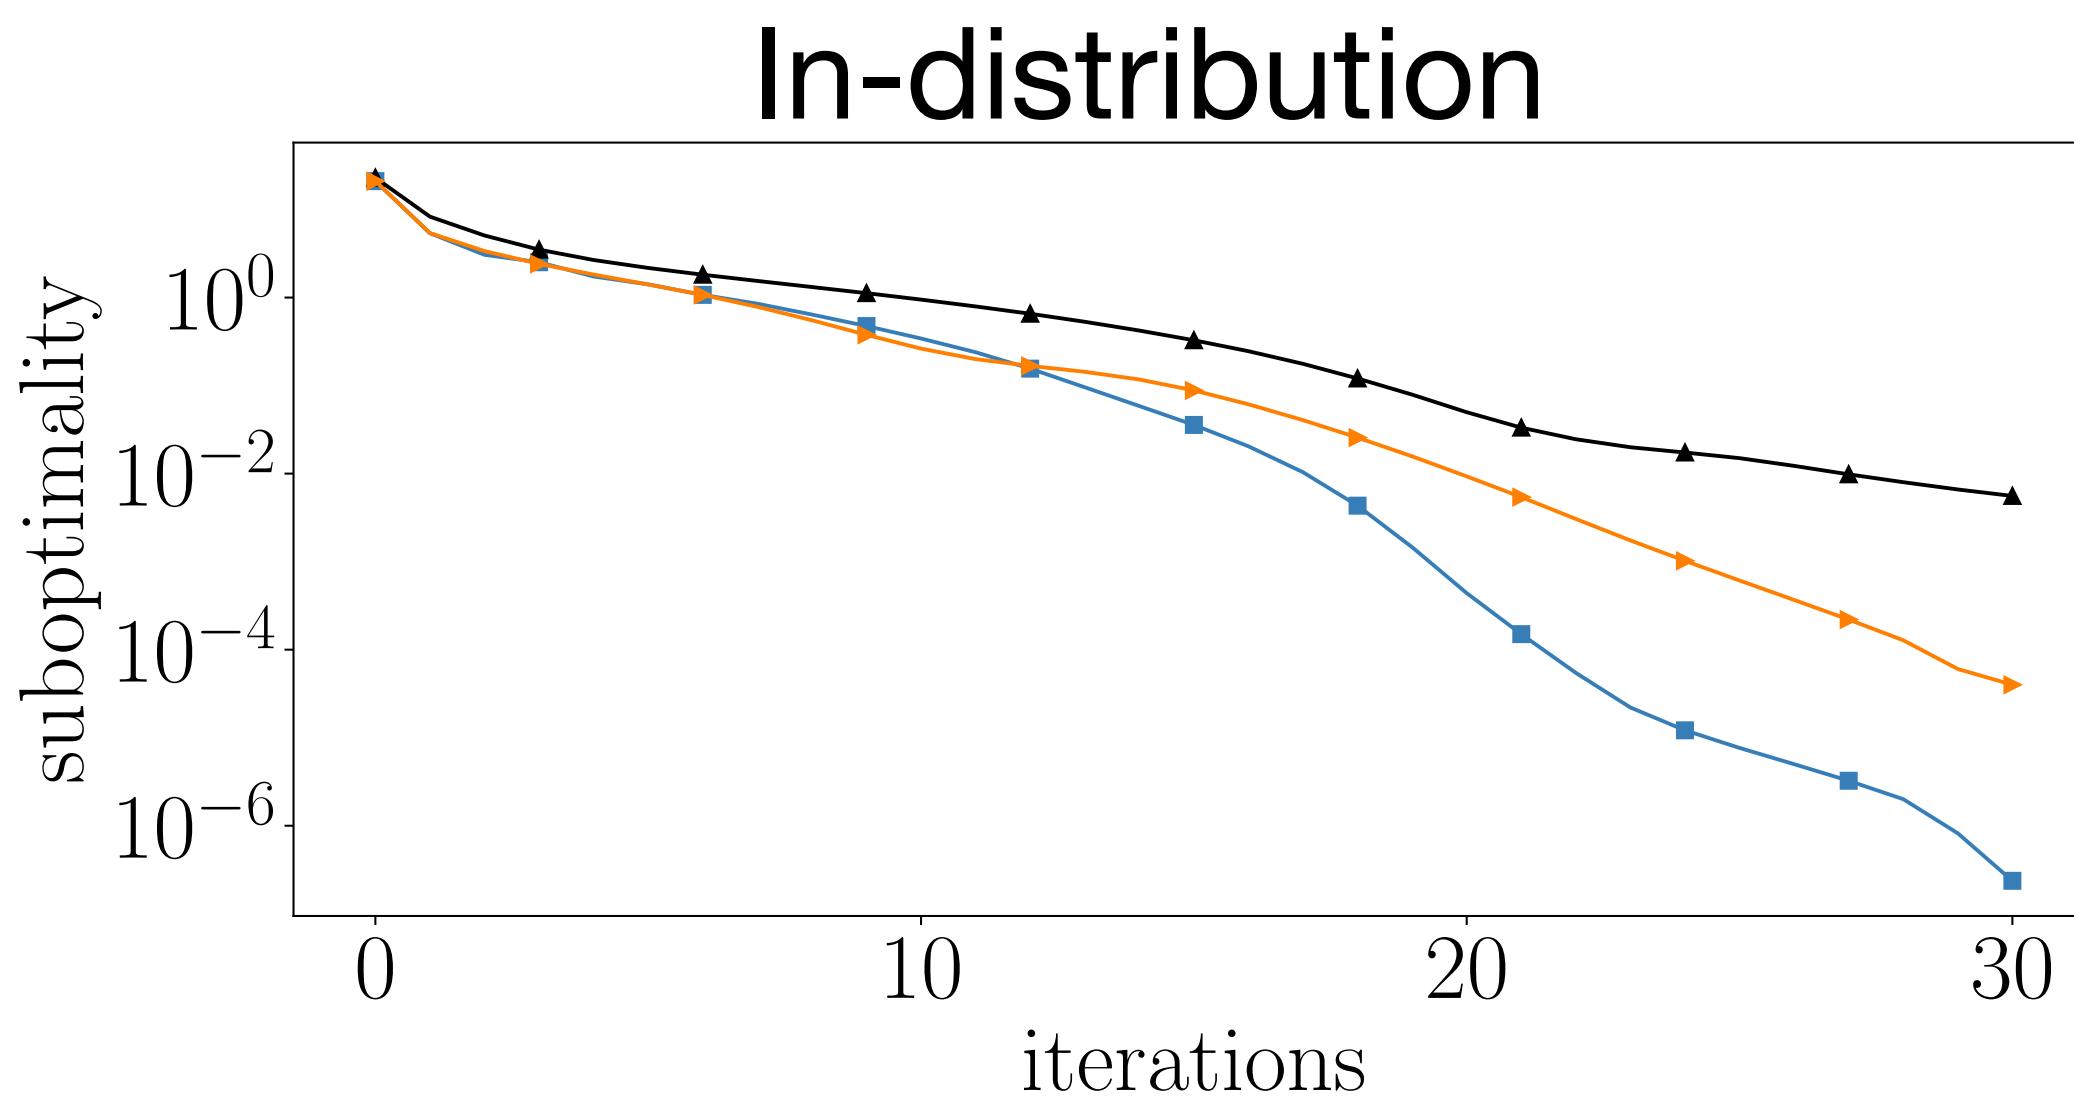
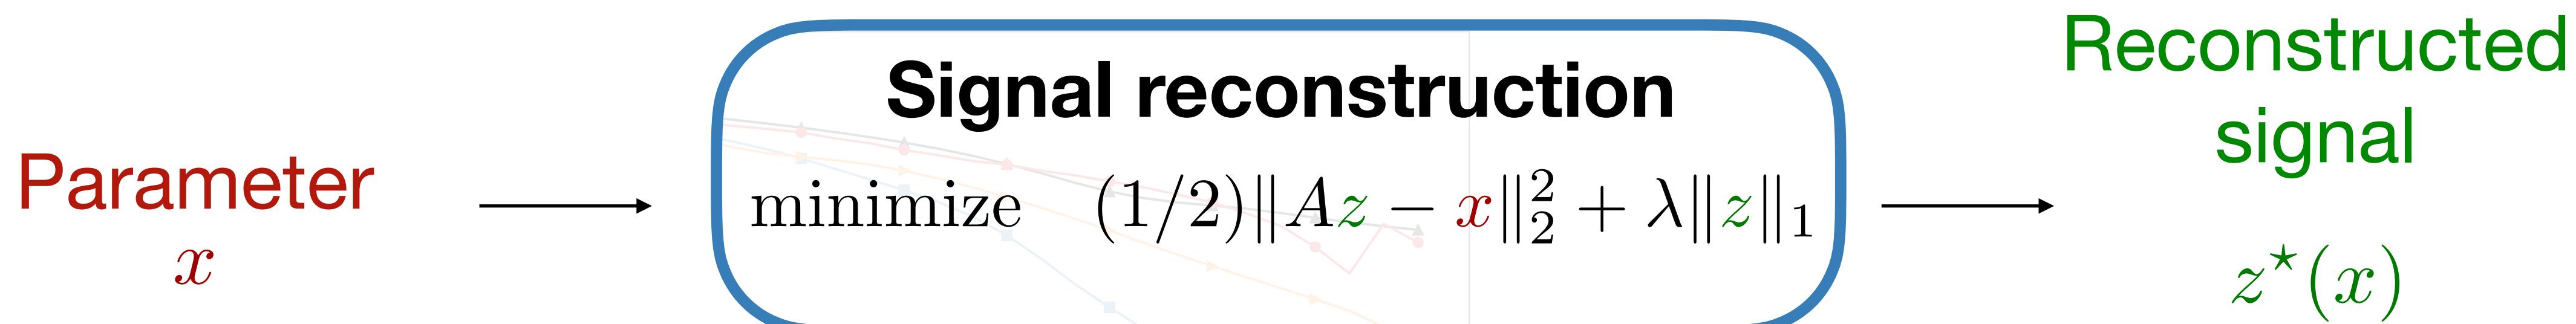
minimize     $(1/N) \sum_{i=1}^N \ell(z_\theta^K(x_i), x_i) + \lambda((\gamma(\theta) - \gamma^{\text{target}})_+)^2$  ← **Penalty formulation**  
subject to     $y_\theta^{k+1}(x_i) = \text{prox}_{\alpha^k}(z_\theta^k(x_i) - \alpha^k \nabla f(z_\theta^k(x_i), x_i))$   
                   $z_\theta^{k+1}(x_i) = y_\theta^{k+1}(x_i) + \beta^k(y_\theta^{k+1}(x_i) - y_\theta^k(x_i))$   
                   $z_\theta^0(x_i) = 0, y_\theta^0(x_i) = 0$

differentiable optimization  
to compute     $\frac{\partial \gamma(\theta)}{\partial \theta}$



Amos et. al 2017, Agrawal et. al 2019

# Learning robust hyperparameters for sparse coding



We can train and maintain robustness

Guarantee holds for any  $x \in \mathbf{R}^d$

# **Learning hyperparameters for the alternating direction method of multipliers (ADMM)**

# We learn hyperparameters for accelerated ADMM also

Two popular ADMM-based solvers



Stellato et al. 2020



Conic problems

$$\begin{aligned} \min \quad & (1/2)w^T Pw + c^T w \\ \text{s.t.} \quad & Aw + s = b \\ & s \in \mathcal{K} \end{aligned}$$

Convex cone

$$\text{with } x = (P, A, c, b)$$

Accelerated Splitting Conic Solver

$$\text{solve } \begin{bmatrix} P + \sigma I & A^T \\ -A & \rho I \end{bmatrix} \tilde{u}^{k+1} = z^k - \begin{bmatrix} c \\ b \end{bmatrix}$$

$$u^{k+1} = \Pi_{\mathbf{R}^q \times \mathcal{K}^*}(2\tilde{u}^{k+1} - z^k)$$

$$y^{k+1} = z^k + \alpha^k(u^{k+1} - \tilde{u}^{k+1})$$

$$z^{k+1} = y^{k+1} + \beta^k(y^{k+1} - y^k)$$

Time-varying hyperparameters  $(\alpha^k, \beta^k)$

Time-invariant hyperparameters  $(\sigma, \rho)$

Why time-invariant?

1. Amenable to PEP
2. Computational advantages—reuse matrix factorization

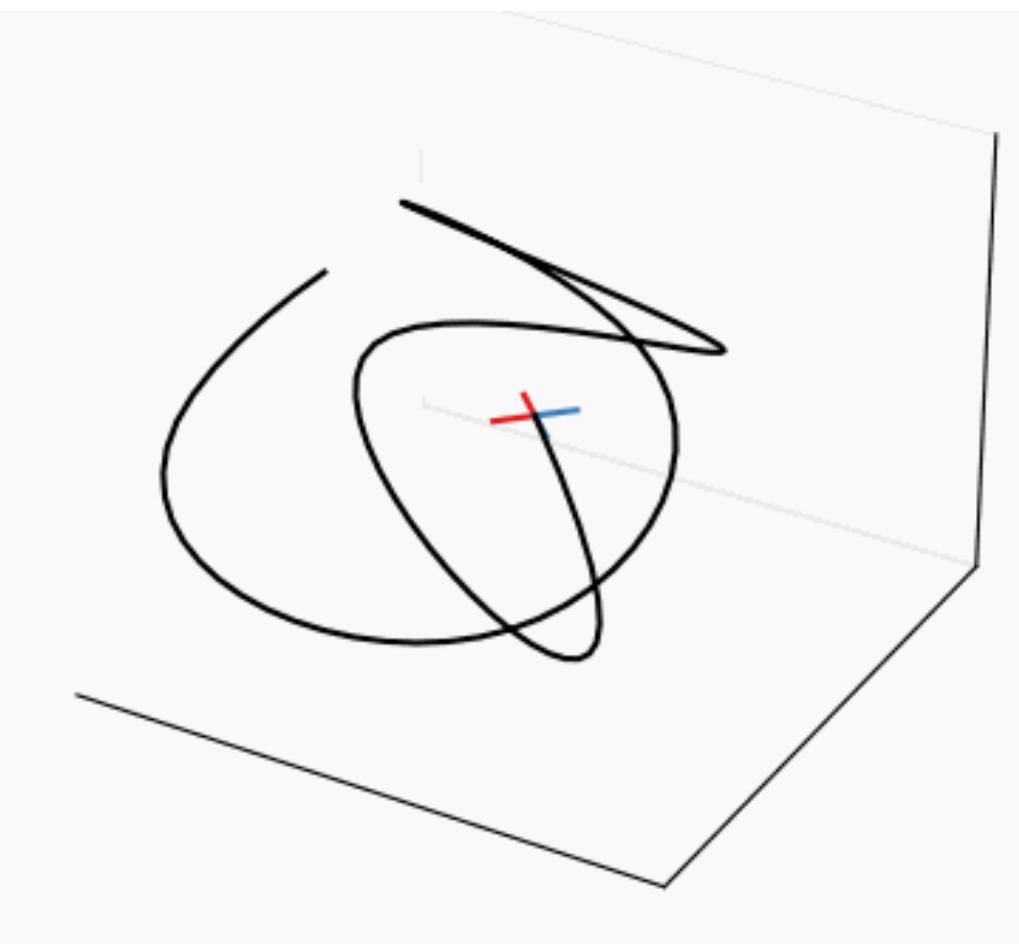
# Model predictive control of a quadcopter

Current state,  
reference trajectory

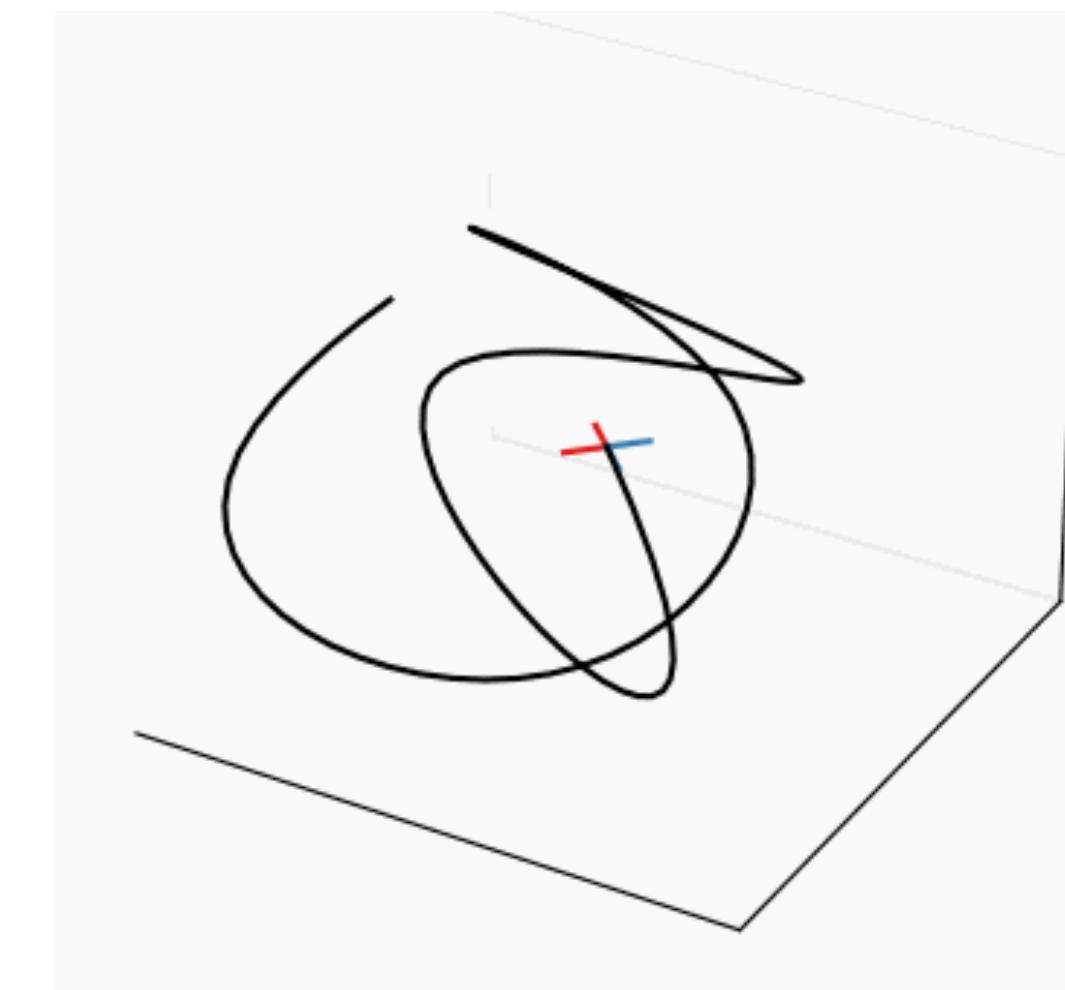
## Quadratic program

$$\begin{aligned} & \text{minimize} && \sum_{t=1}^T \|s_t - s_t^{\text{ref}}\|_2^2 \\ & \text{subject to} && s_{t+1} = As_t + Bu_t \\ & && s_t \in \mathcal{S}, \quad u_t \in \mathcal{U} \\ & && s_0 = s_{\text{init}} \end{aligned}$$

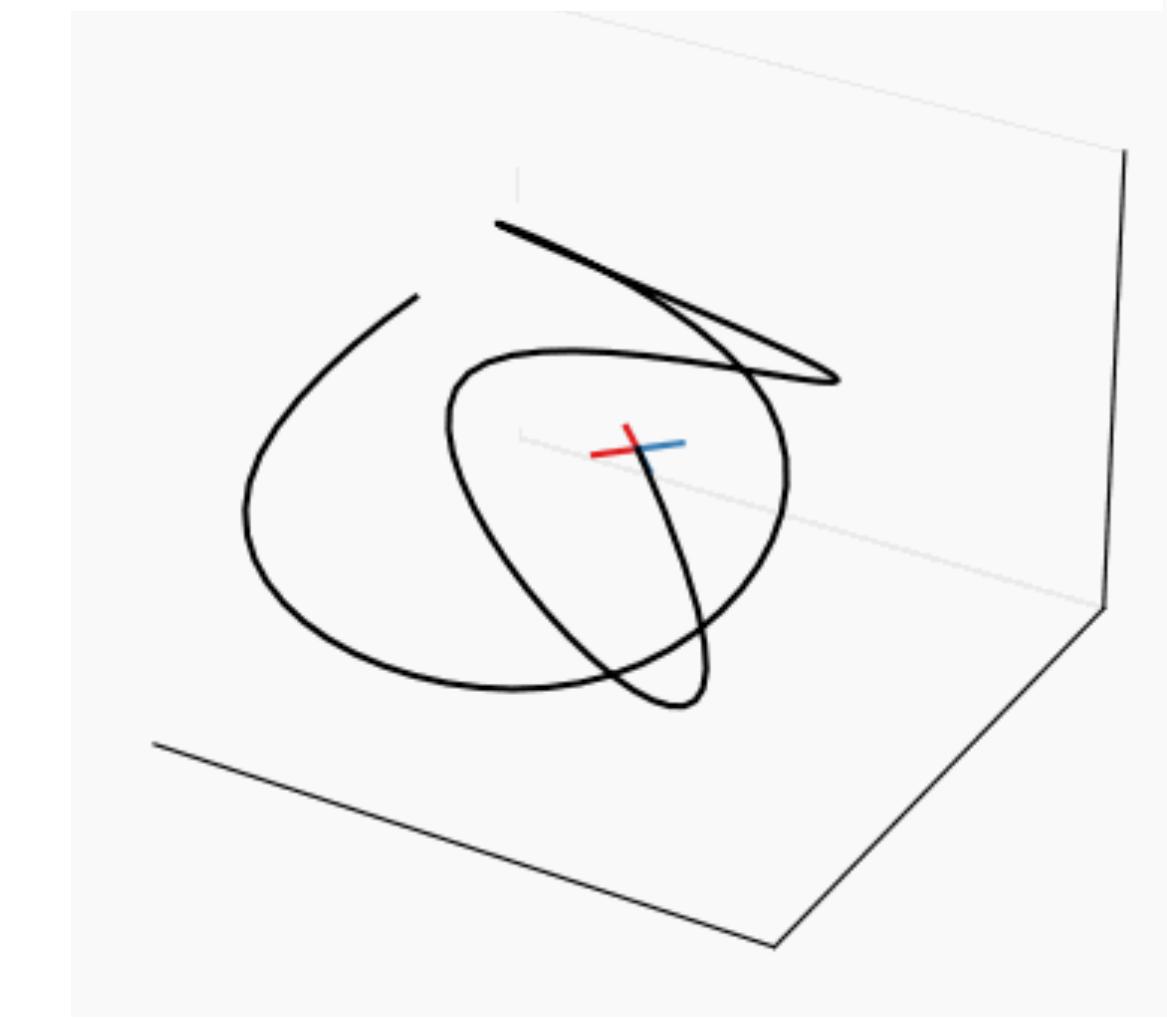
Control  
inputs



Nearest neighbor  
80 iterations



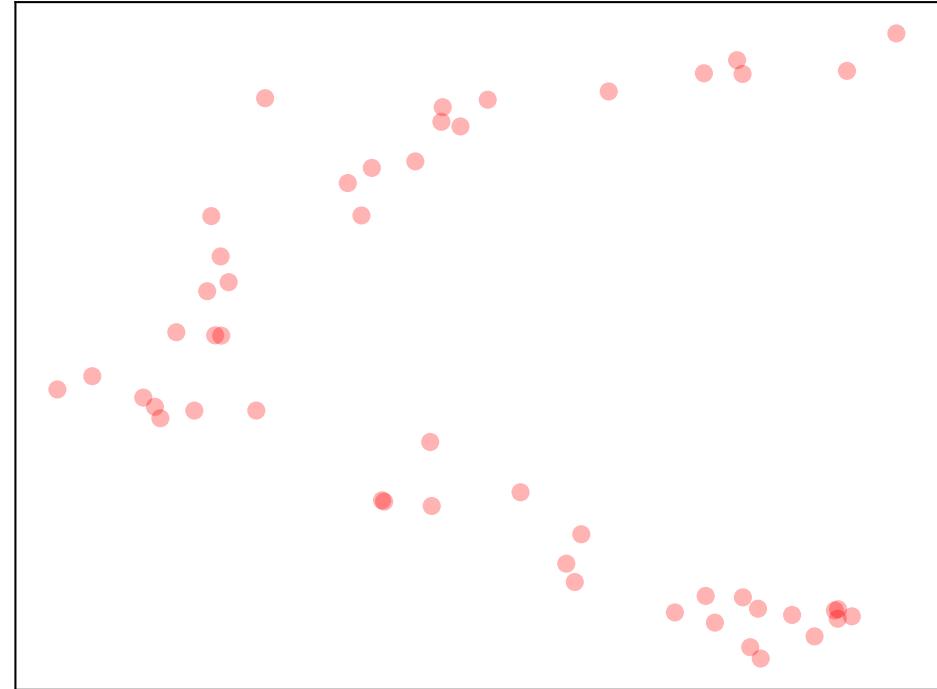
Previous solution  
80 iterations



Learned accel + robust  
20 iterations

With learning, we  
can track the  
trajectory well

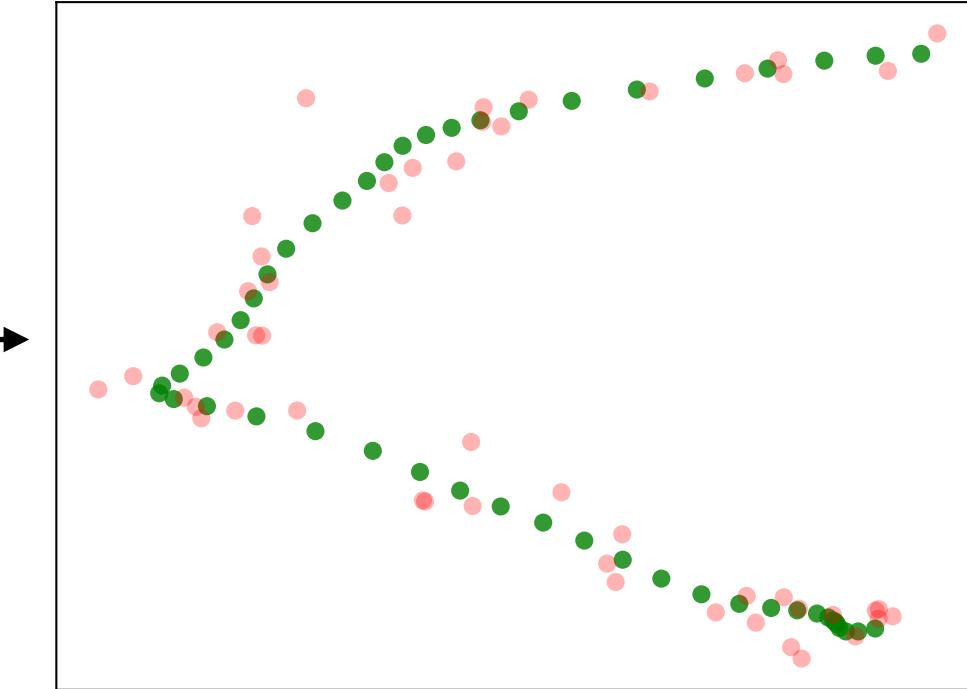
# Robust Kalman filtering



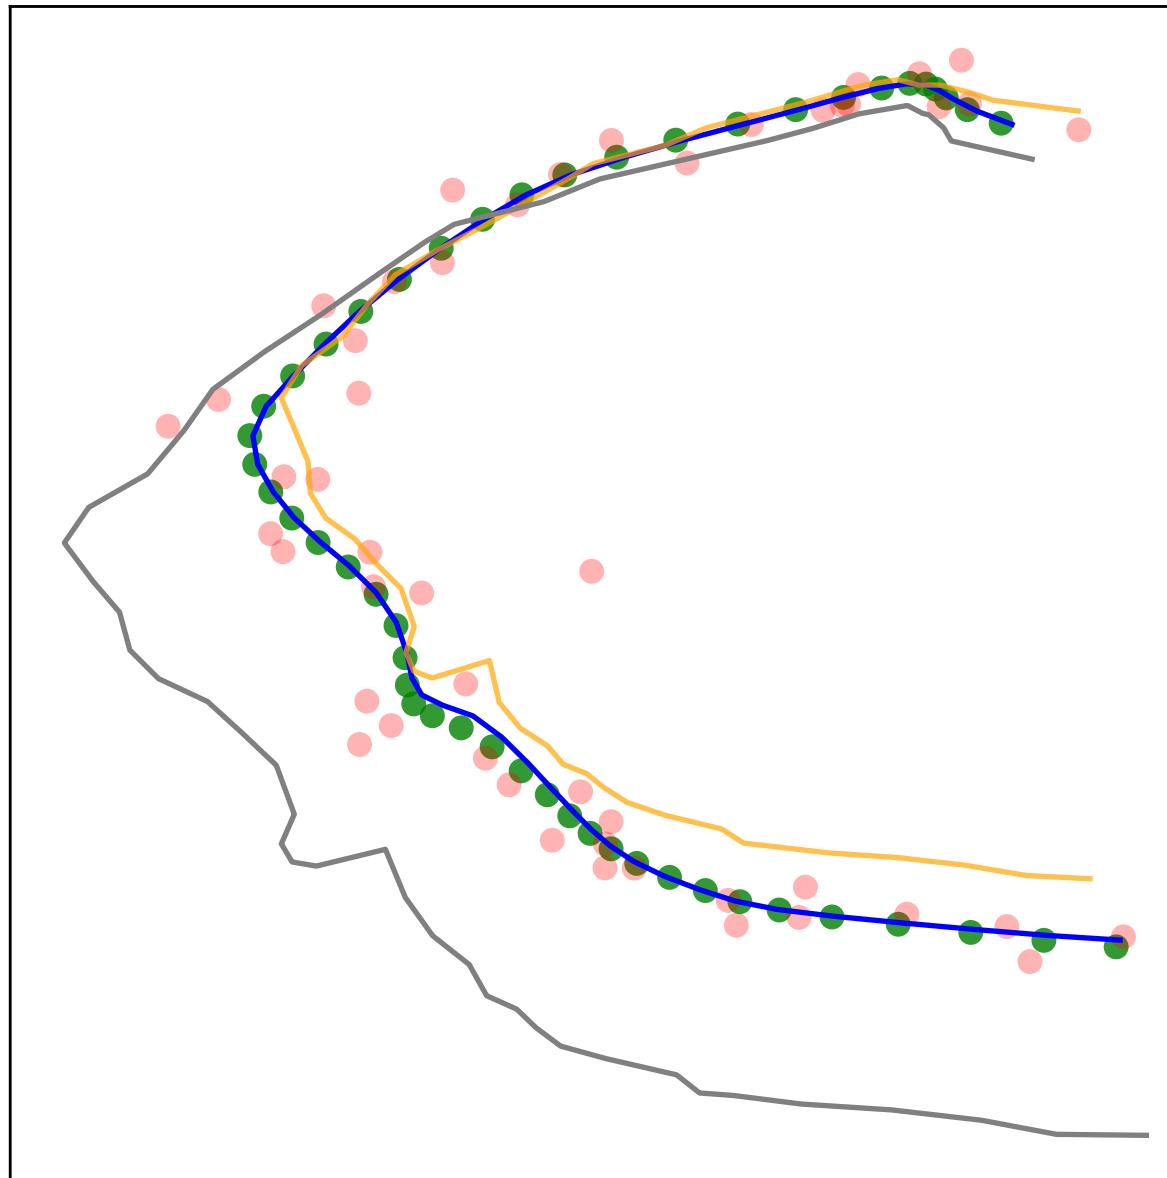
## Second-order cone program

minimize       $\sum_{t=0}^{T-1} \|w_t\|_2^2 + \psi_\rho(v_t)$   
subject to     $s_{t+1} = As_t + Bu_t$   
                   $y_t = Cs_t + v_t$

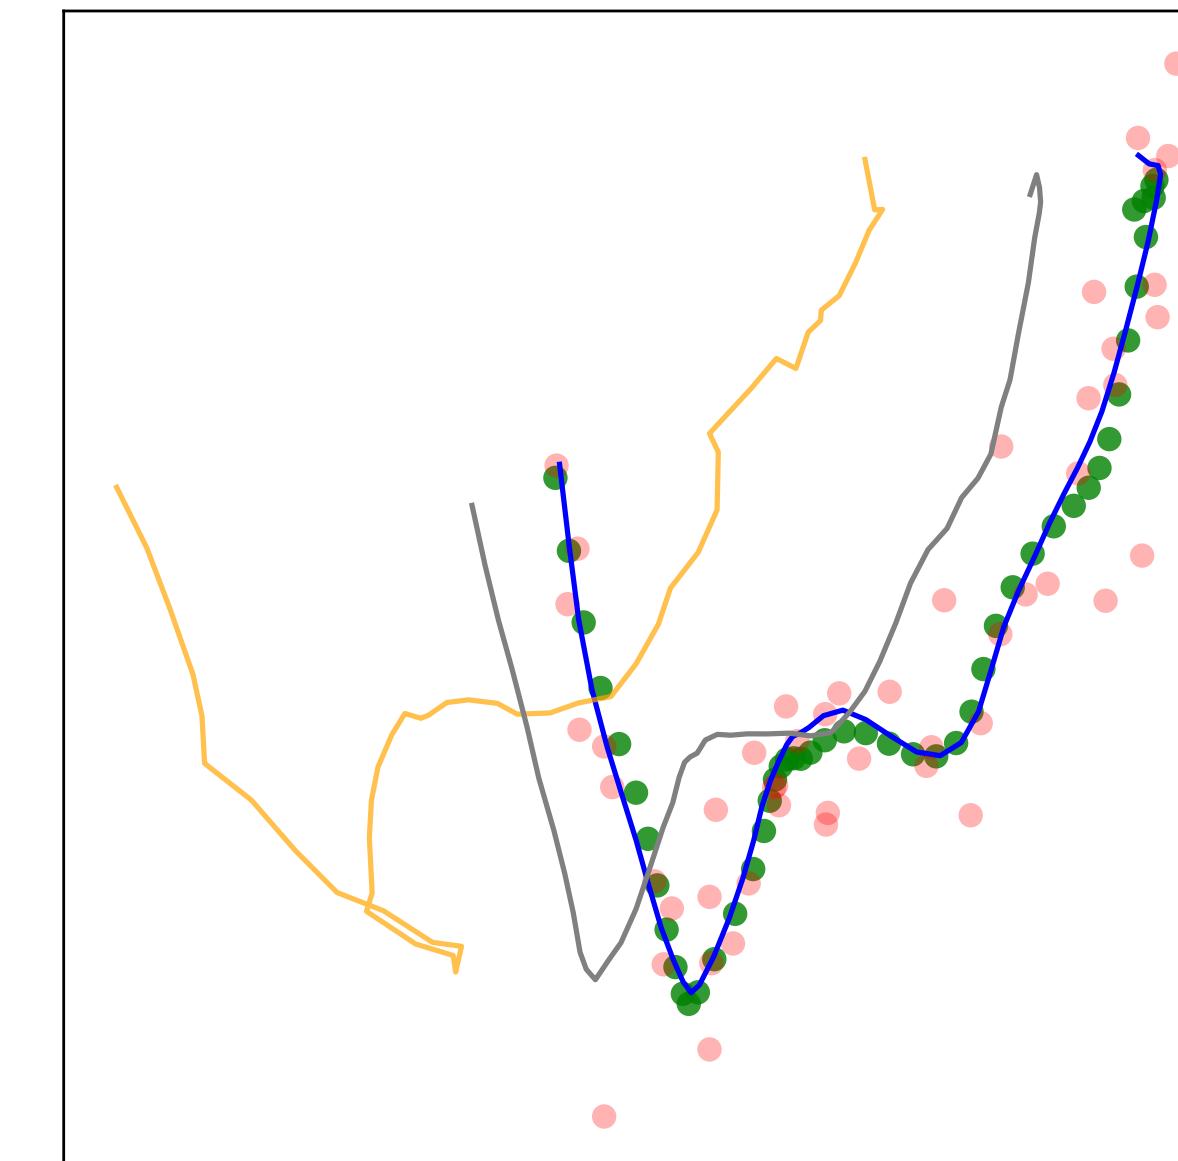
Huber loss



In-distribution



Out-of-distribution



5 iterations

No learning

Learned hyperparameters

Learned acceleration + Robust

Learning acceleration algorithms w/  
robustness tracks the optimal solution

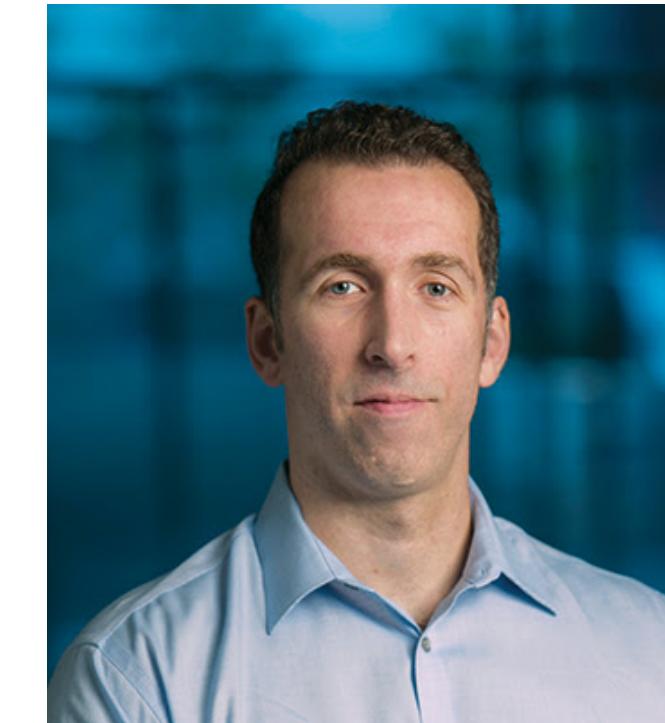
# Acknowledgements



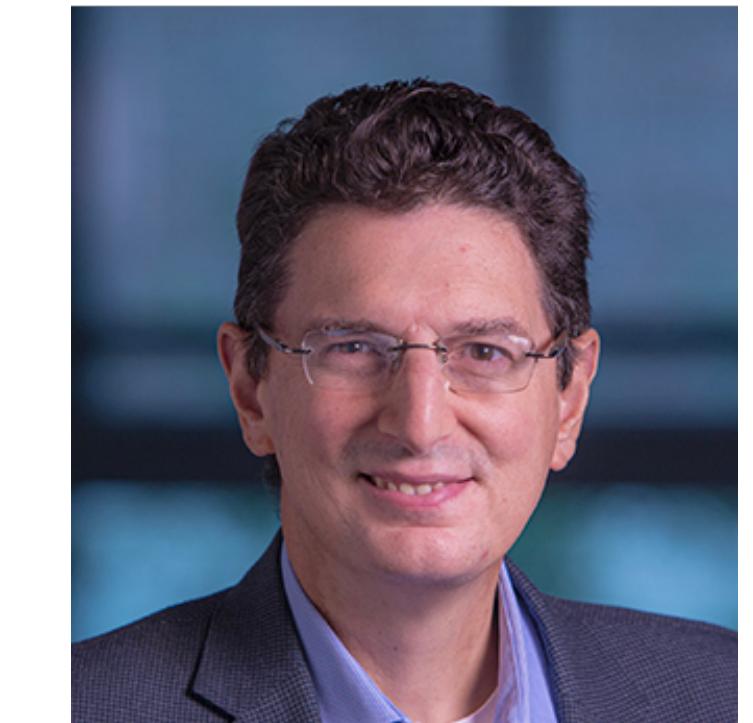
Bartolomeo Stellato



Jinho Bok



Nikolai Matni



George Pappas



**Learning Algorithm Hyperparameters for Fast Parametric Convex Optimization**

R. Sambharya, B. Stellato

<https://arxiv.org/pdf/2411.15717>

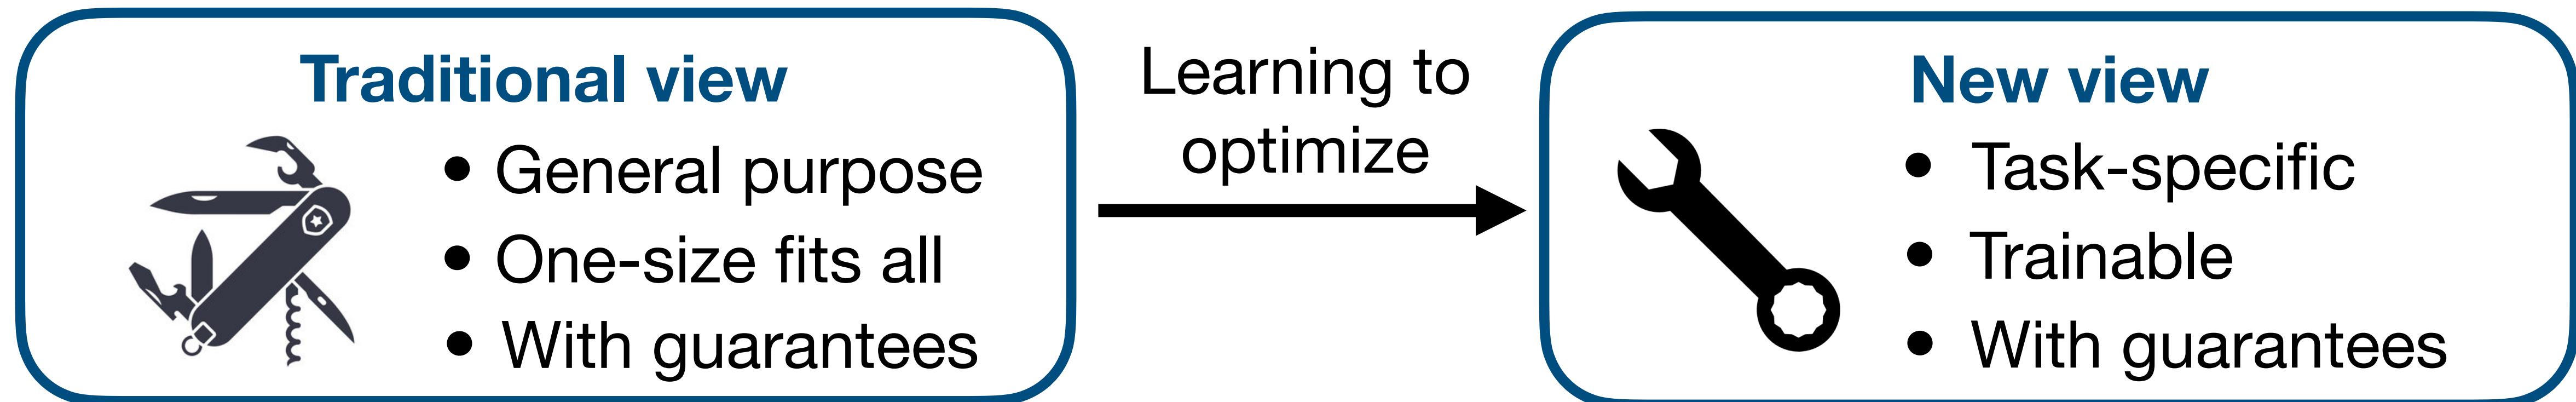


**Learning Acceleration Algorithms for Fast Parametric Convex Optimization with Certified Robustness**

R. Sambharya, J. Bok, N. Matni, G. Pappas

<https://arxiv.org/pdf/2507.16264>

# Conclusion



Takeaways from this talk specifically

- Only **learning** the **hyperparameter sequence** dramatically improves performance
- Very **low amount of training data** needed
- We evaluate and train for **robustness** using PEP



[sambhar9@seas.upenn.edu](mailto:sambhar9@seas.upenn.edu)



[rajivsambharya.github.io](https://rajivsambharya.github.io)